

# Part II

# Processes and Threads

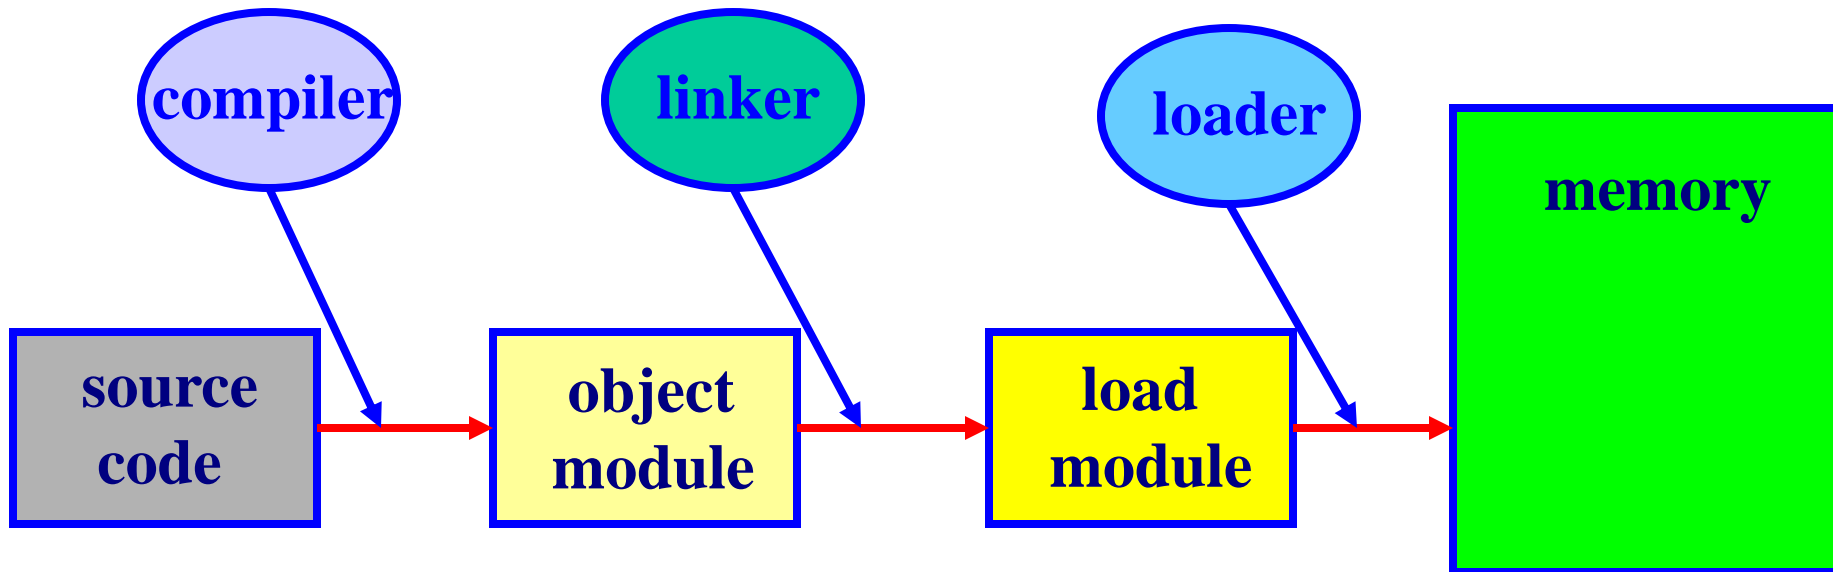
## Process Basics

*Program testing can be used to show the presence of bugs,  
but never to show their absence*

1

# ***From Compilation to Execution***

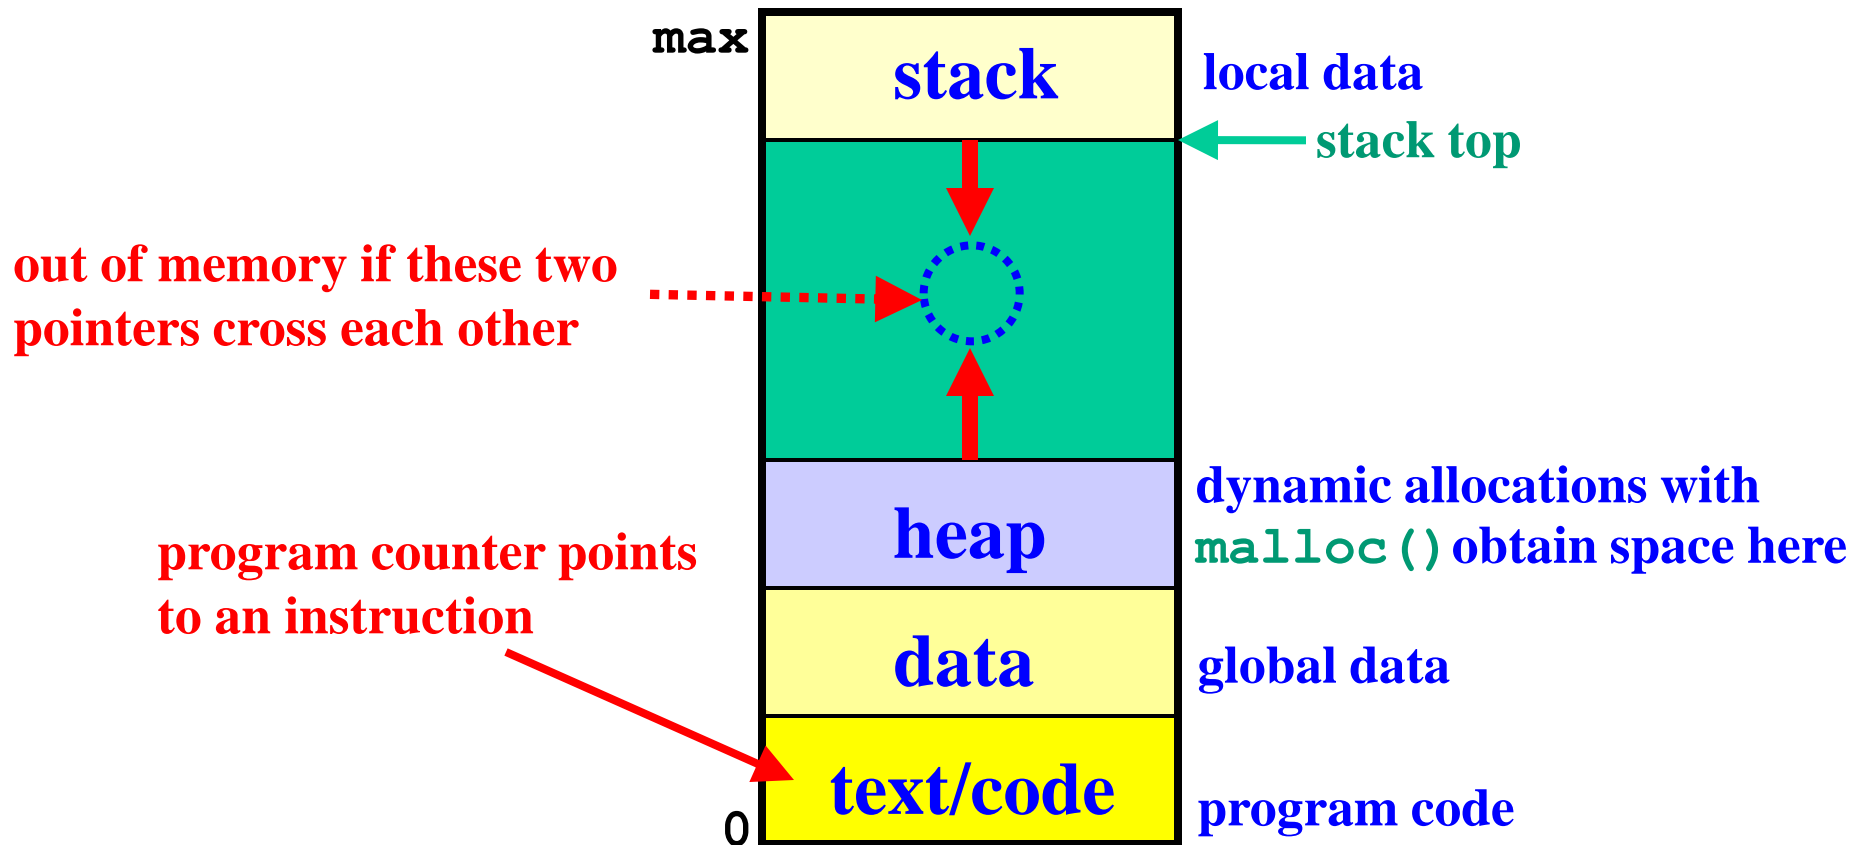
- A compiler compiles source files to `.o` files.
- A linker links `.o` files and other libraries together, producing a binary executable (e.g., `a.out`).
- A loader loads a binary executable into memory for execution.



# ***What Is a Process?***

- When the OS runs a program (i.e., a binary executable), this program is loaded into memory and the control is transferred to this program's first instruction. Then, the program runs.
- ***A process is a program in execution.***
- A process is more than a program, because a process has a **program counter**, **stack**, **data section**, **code section**, etc (i.e., the runtime stuffs).
- Moreover, multiple processes may be associated with one program (e.g., run the same program, say **a.out**, multiple times at the same time).

# Process Space



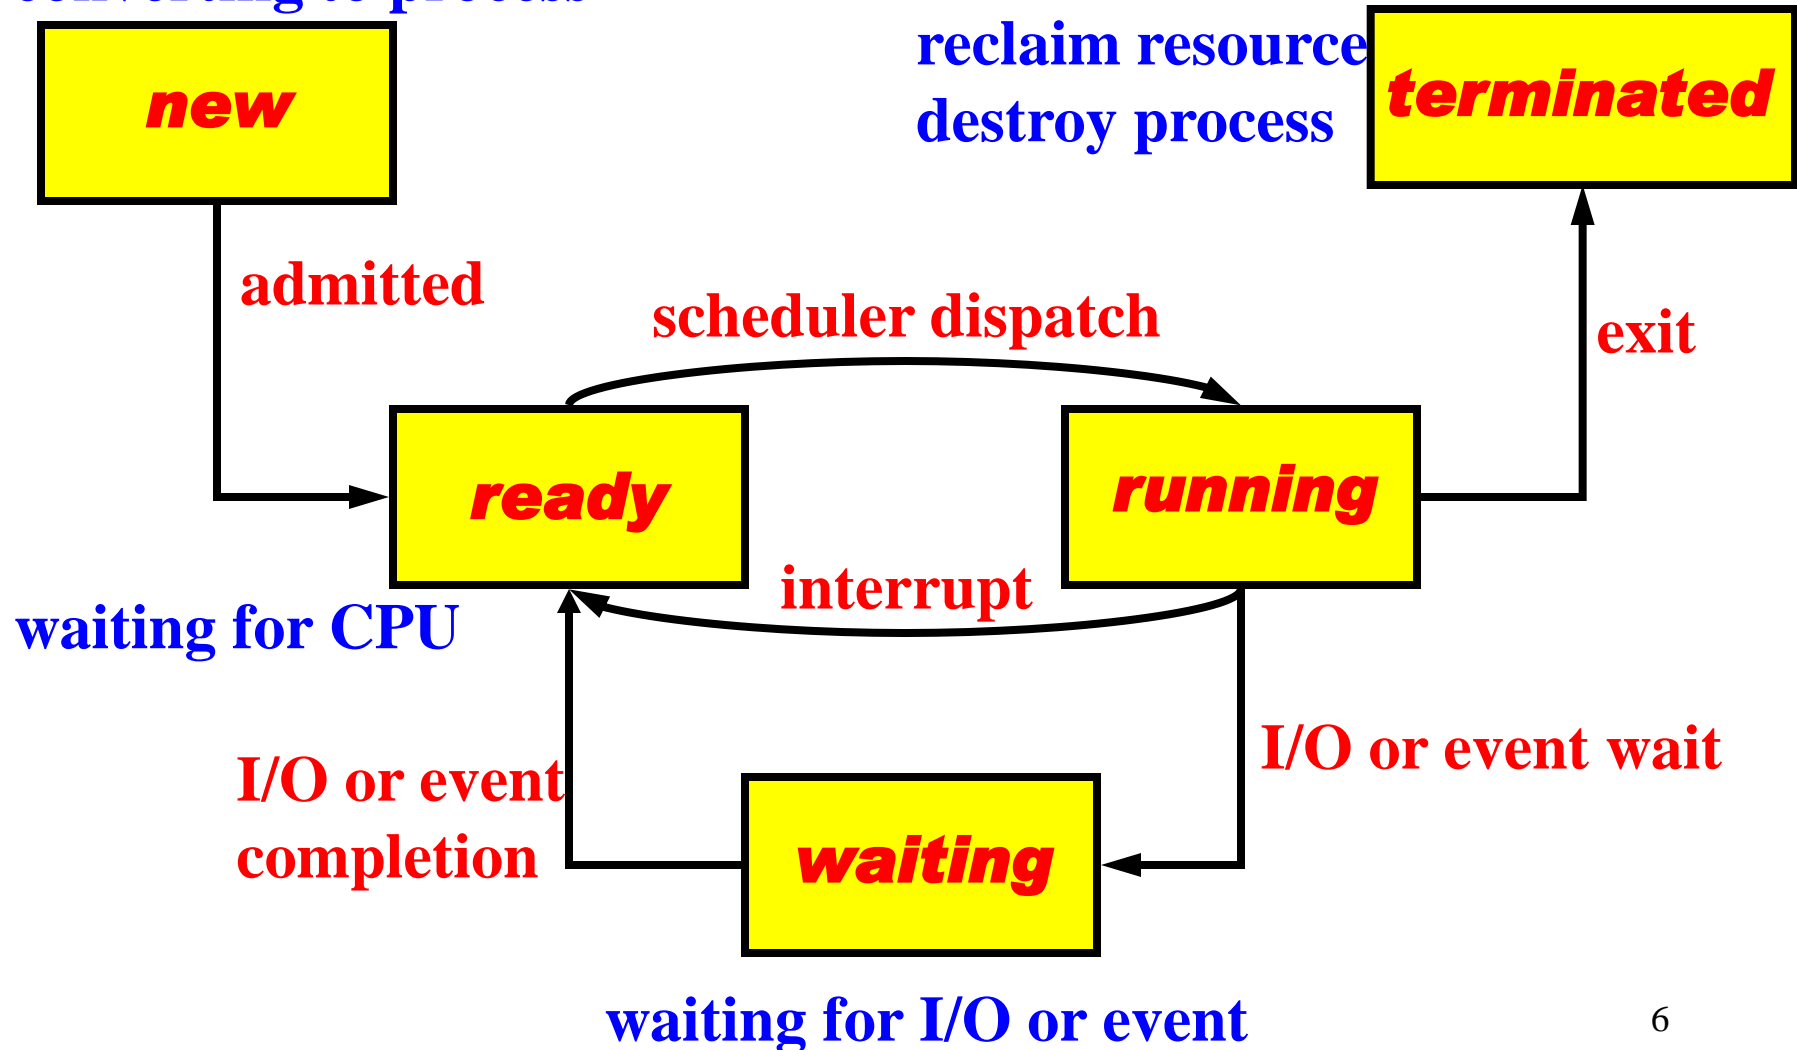
# ***Process States***

At any moment, a process can be in one of the five states: **new**, **running**, **waiting**, **ready** and **terminated**.

- ***New***: The process is being created
- ***Running***: The process is executing on a CPU
- ***Waiting***: The process is waiting for some event to occur (e.g., waiting for I/O completion)
- ***Ready***: The process has everything but the CPU. It is waiting to be assigned to a processor.
- ***Terminated***: The process has finished execution.

# ***Process State Diagram***

converting to process



# Process Representation in OS

pointer	process state
process ID	
program counter	
registers	
scheduling info	
memory limits	
list of open files	
⋮	

- Each process is assigned a unique number, the **process ID**.
- Process info are stored in a table, the **process control block (PCB)**.
- These PCBs are chained into a number of lists. For example, all processes in the ready state are in the **ready queue**.

# ***Process Scheduling: 1/2***

- Since the number of processes may be larger than the number of available CPUs, the OS must maintain **maximum CPU utilization**.
- To determine which process can do what, processes are chained into a number of **scheduling queues**.
- For example, in addition to the ready queue, each event may have its own scheduling queue (*i.e.*, waiting queue).



# ***Process Scheduling: 2/2***

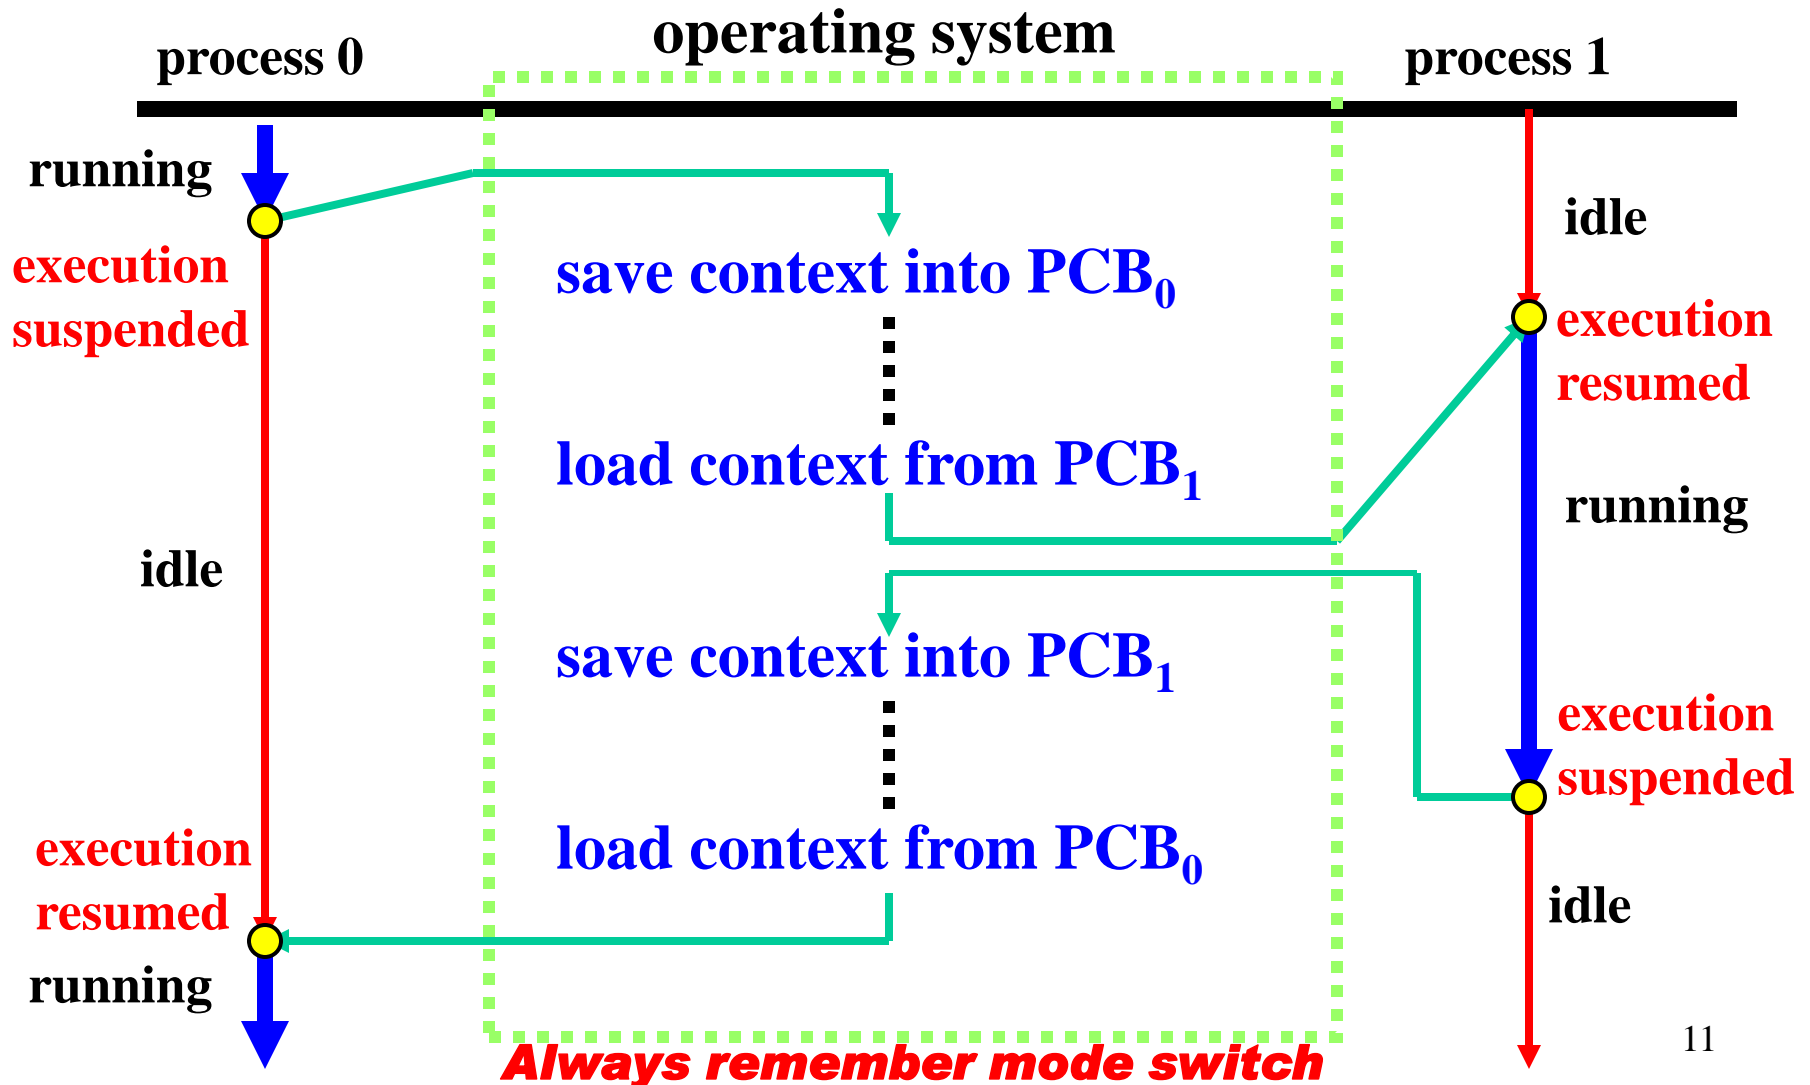
- The ready queue, which may be organized into several sub-queues, has all processes ready to run.
- The OS has a ***CPU scheduler***.
- When a CPU is free, the CPU scheduler looks at the ready queue, picks a process, and resumes it.
- The way of picking a process from the ready queue is referred to as ***scheduling policy***.
- Scheduling policy is unimportant to this course because we ***cannot make any assumption*** about it.

# Context Switch: 1/2

- **What is a process context?** The **context** of a process includes process ID, process state, the values of CPU registers, the program counter, and other memory/file management information (i.e., execution environment).
- **What is a context switch?** After the CPU scheduler selects a process and before allocates CPU to it, the CPU scheduler must
  - save the context of the currently running process,
  - put it back to the ready queue or waiting state,
  - load the context of the selected process, and
  - go to the saved program counter.

**mode switching may be needed**

# Context Switch: 2/2



# ***Operations on Processes***

- There are three commonly seen operations:
  - ❖ **Process Creation:** Create a new process. The newly created is the child of the original. Unix uses `fork()` to create new processes.
  - ❖ **Process Termination:** Terminate the execution of a process. Unix uses `exit()`.
  - ❖ **Process Join:** Wait for the completion of a child process. Unix uses `wait()`.
- `fork()`, `exit()` and `wait()` are system calls.

# ***Some Required Header Files***

- Before you use processes, include header files `sys/types.h` and `unistd.h`.
- `sys/types.h` has all system data types, and `unistd.h` declares standard symbolic constants and types.

```
#include <sys/types.h>
#include <unistd.h>
```

# ***The `fork()` System Call***

- The purpose of `fork()` is to create a child process. The creating and created processes are the **parent** and **child**, respectively.
- `fork()` does not require any argument!
- If the call to `fork()` is successful, Unix creates an **identical** but **separate** address space for the child process to run.
- Both processes start running with the instruction following the `fork()` system call.


# ***fork()*** ***Return Values***

- A negative value means the creation of a child process was unsuccessful.
- A zero means the process is a child.
- Otherwise, `fork()` returns the process ID of the child process. The ID is of type `pid_t`.
- Function `getpid()` returns the process ID of the caller.
- Function `getppid()` returns the parent's process ID. If the calling process has no parent, `getppid()` returns 1.

# ***Before the Execution of `fork()`***

parent

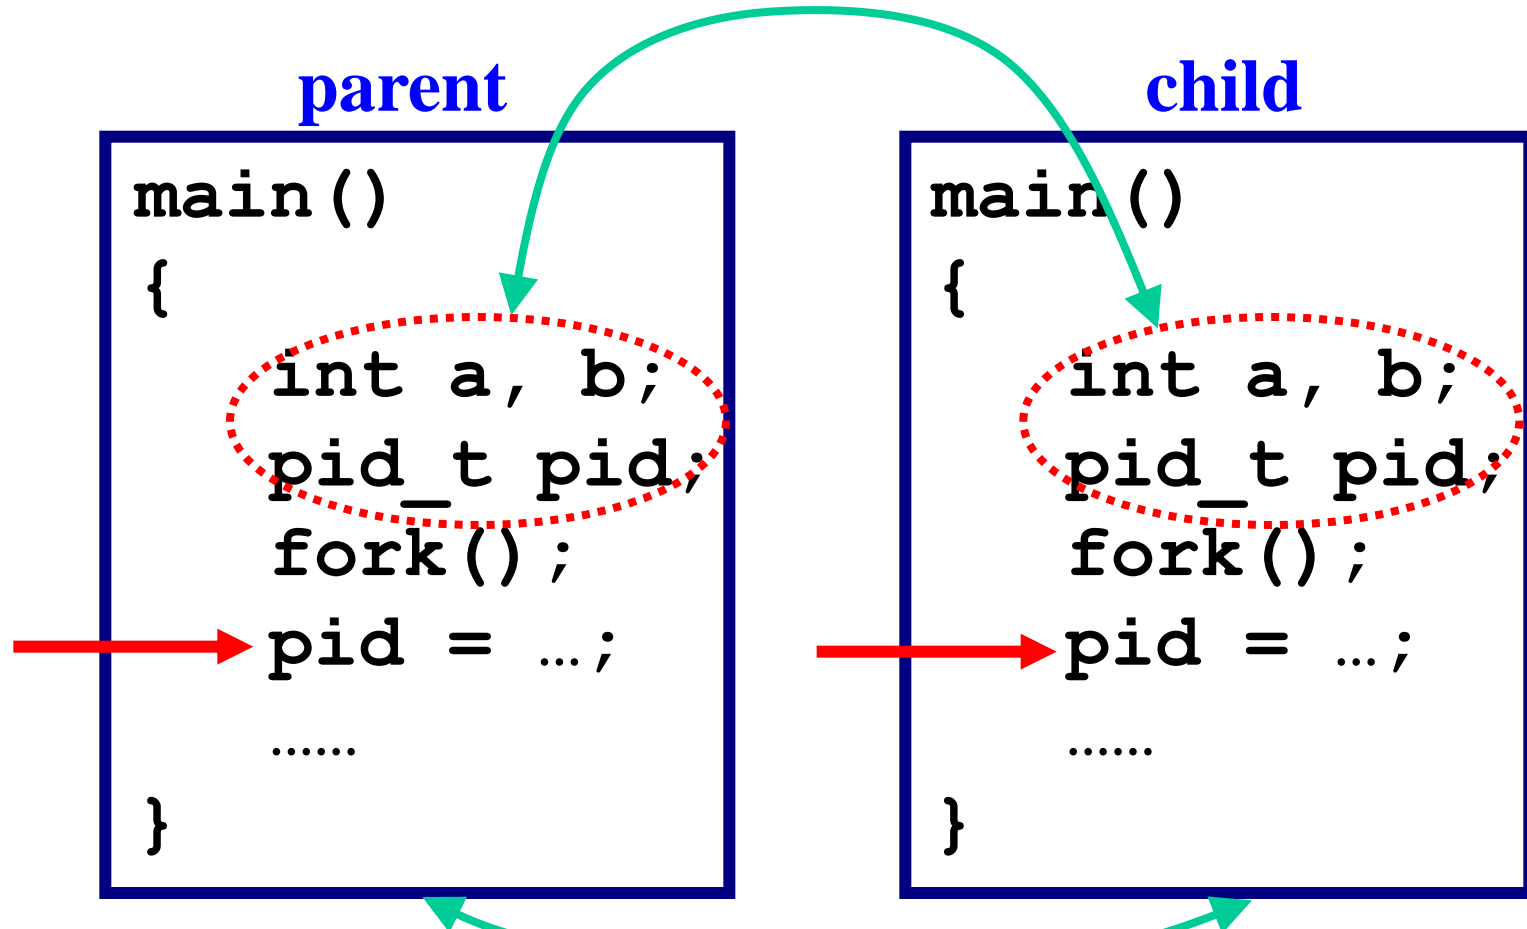
```
main()  
{  
    int a, b;  
    pid_t pid  
    fork();  
    pid = ...;  
    .....  
}
```





# ***After the Execution of `fork()`***

in different address spaces



two independent and separate address spaces


# Example 1: 1/2

**fork-1.c**

```
#include <stdio.h>
#include <string.h>
#include <sys/types.h>
#include <unistd.h>

void main(void)
{
    pid_t  MyID, ParentID;
    int    i;
    char   buf[100];

    fork(); // create a child process
    MyID    = getpid(); // get my process ID
    ParentID = getppid(); // get my parent's process ID
    for (i = 1; i <= 200; i++) {
        sprintf(buf, "From MyID=%ld, ParentID=%ld, i=%3d\n",
                    MyID, ParentID, i);
        write(1, buf, strlen(buf)); // why don't we
    }                                // use printf?
}
```

 this is stdout

# Example 1: 2/2

```
hilbert.cs.mtu.edu - PuTTY
From MyID = 19087, ParentID = 19004, i = 193
From MyID = 19088, ParentID = 19087, i = 88
From MyID = 19087, ParentID = 19004, i = 194
From MyID = 19088, ParentID = 19087, i = 89
From MyID = 19087, ParentID = 19004, i = 195
From MyID = 19088, ParentID = 19087, i = 90
From MyID = 19087, ParentID = 19004, i = 196
From MyID = 19088, ParentID = 19087, i = 91
From MyID = 19087, ParentID = 19004, i = 197
From MyID = 19088, ParentID = 19087, i = 92
From MyID = 19087, ParentID = 19004, i = 198
From MyID = 19088, ParentID = 19087, i = 93
From MyID = 19087, ParentID = 19004, i = 199
From MyID = 19088, ParentID = 19087, i = 94
From MyID = 19087, ParentID = 19004, i = 200
From MyID = 19088, ParentID = 19087, i = 95
From MyID = 19088, ParentID = 19087, i = 96
From MyID = 19088, ParentID = 19087, i = 97
From MyID = 19088, ParentID = 19087, i = 98
From MyID = 19088, ParentID = 19087, i = 99
From MyID = 19088, ParentID = 19087, i = 100
From MyID = 19088, ParentID = 19087, i = 101
From MyID = 19088, ParentID = 19087, i = 102
From MyID = 19088, ParentID = 19087, i = 103
```

Processes 19087 and 19088 run concurrently

Parent: 19087  
Child : 19088

19004  
↓  
19087  
↓  
19088

Parent 19087's parent is 19004, the shell that executes `fork-1`

# *fork(): A Typical Use*

```
main(void)
{
    pid_t pid;

    pid = fork();
    if (pid < 0)
        printf("Oops!");
    else if (pid == 0)
        child();
    else // pid > 0
        parent();
}
```

```
void child(void)
{
    int i;
    for (i=1; i<=10; i++)
        printf(" Child:%d\n", i);
    printf("Child done\n");
}

void parent(void)
{
    int i;
    for (i=1; i<=10; i++)
        printf("Parent:%d\n", i);
    printf("Parent done\n");
}
```

we use `printfs` here to save space.

# ***Before the Execution of fork()***

parent

```
main(void)  pid = ?
{
  → pid = fork();
    if (pid == 0)
      child();
    else
      parent();
}

void child(void)
{ ..... }

void parent(void)
{ ..... }
```

# After the Execution of `fork()` 1/2

parent

```
main(void) pid=123
```

```
{  
    pid = fork();  
    → if (pid == 0)  
        child();  
    else  
        parent();  
}
```

```
void child(void)  
{ ..... }
```

```
void parent(void)  
{ ..... }
```

child

```
main(void) pid=0
```

```
{  
    pid = fork();  
    → if (pid == 0)  
        child();  
    else  
        parent();  
}
```

```
void child(void)  
{ ..... }
```

```
void parent(void)  
{ ..... }
```

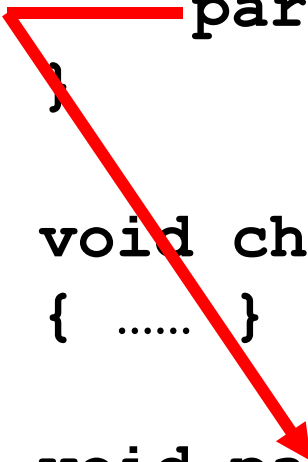
in two different address spaces

# After the Execution of `fork()` 2/2

parent

```
main(void) pid=123
{
    pid = fork();
    if (pid == 0)
        child();
    else
        parent();
}

void child(void)
{ ..... }
```

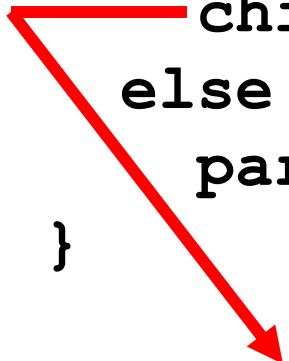


```
void parent(void)
{ ..... }
```

child

```
main(void) pid=0
{
    pid = fork();
    if (pid == 0)
        child();
    else
        parent();
}

void child(void)
{ ..... }
```



```
void parent(void)
{ ..... }
```

# Example 2: 1/2

```
#include .....           // child exits first           fork-2.c
void main(void)
{
    pid_t  pid;

    pid = fork();
    if (pid == 0) {           // child here
        printf("From child %ld: my parent is %ld\n",
               getpid(), getppid());
        sleep(5);
        printf("From child %ld 5 sec later: parent %ld\n",
               getpid(), getppid());
    }
    else {                   // parent here
        sleep(2);
        printf("From parent %ld: child %ld, parent %ld\n",
               getpid(), pid, getppid());
        printf("From parent: done!\n");
    }
}
```

**force the child to print after the parent terminates**

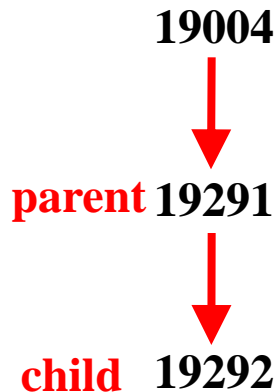
The diagram illustrates the execution flow of the program. A red dotted arrow points from the `sleep(5);` line in the child process to the `sleep(2);` line in the parent process, indicating that the parent waits for the child to finish. Red dotted circles highlight the `sleep(5);` and `sleep(2);` lines, emphasizing the timing of the delays.



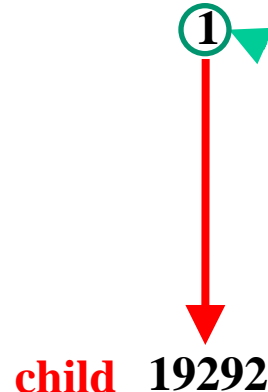
# Example 2: 2/2

```
hilbert.cs.mtu.edu - PuTTY
HILBERT:/home/csdept/shene/CS3331/O3-Process(41) fork-2
  From child 19292: my parent is 19291
From parent 19291: my child is 19292, my parent is 19004
From parent: done!
  From child 19292 5 sec later: my parent is 1
```

while parent is still running



parent terminated



Orphan children have a new parent `init`, the Unix process that spawns all processes

1 is the ID of `init`, the Unix process that spawns all processes

19004 is the shell process that executes 19291

# Example 3: 1/2

```
#include ..... // separate address spaces fork-3.c

void main(void)
{
    pid_t  pid;
    char    out[100];
    int     i = 10, j = 20;

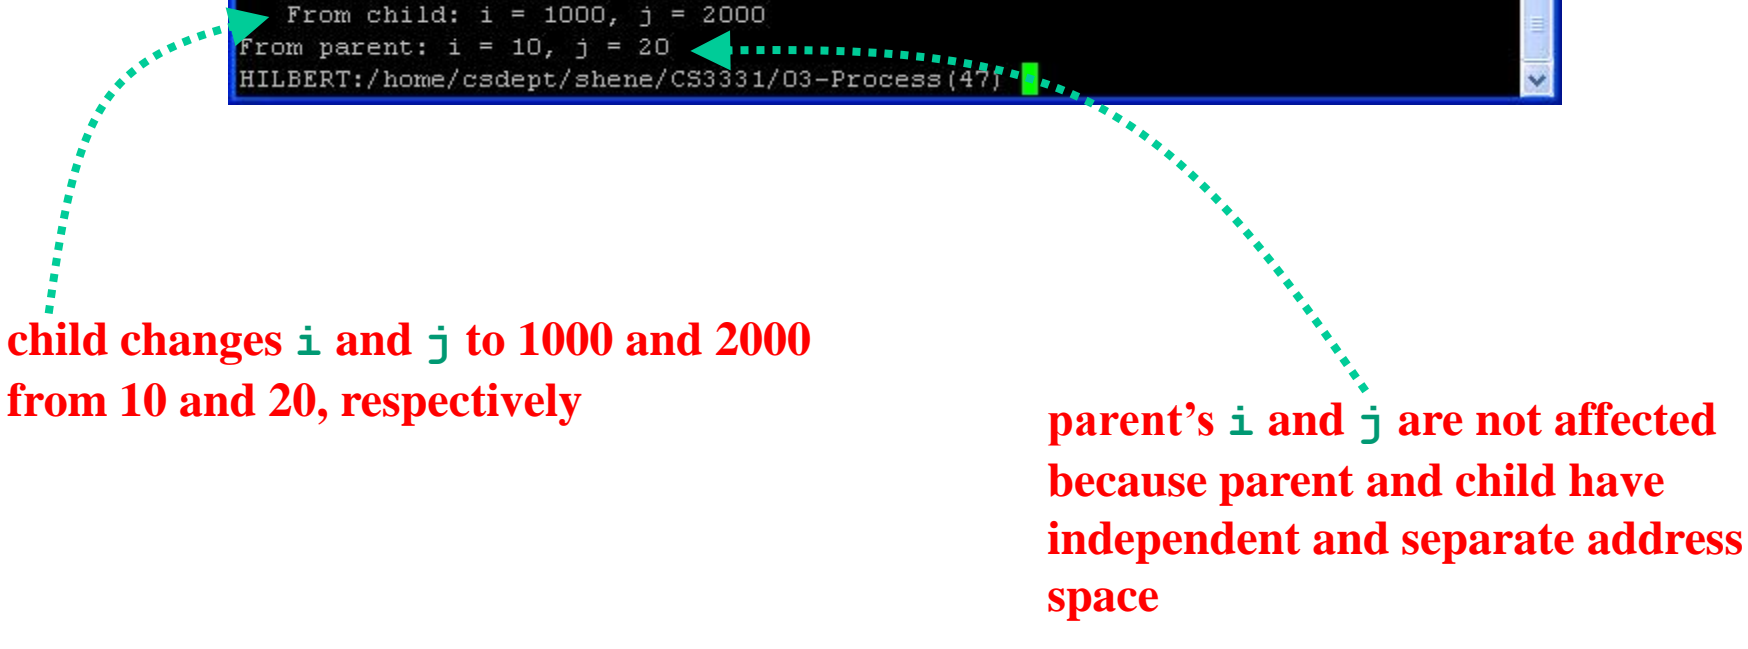
    if ((pid = fork()) == 0) { // child here
        i = 1000; j = 2000;    // child changes values
        sprintf(out, "    From child: i=%d, j=%d\n", i, j);
        write(1, out, strlen(out));
    }
    else { // parent here
        sleep(3);
        sprintf(out, "From parent: i=%d, j=%d\n", i, j);
        write(1, out, strlen(out));
    }
}
```

**this is equivalent to**  
pid = fork();  
if (pid == 0) ...

**force the parent to print after the child terminated**

# Example 3: 2/2

```
HILBERT:/home/csdept/shene/CS3331/O3-Process(46) fork-3
  From child: i = 1000, j = 2000
  from parent: i = 10, j = 20
HILBERT:/home/csdept/shene/CS3331/O3-Process(47)
```



child changes **i** and **j** to 1000 and 2000  
from 10 and 20, respectively

parent's **i** and **j** are not affected  
because parent and child have  
independent and separate address  
space

# ***The wait() System Call***

- The `wait()` system call blocks the caller until ***one of its child processes*** exits or a signal is received.
- `wait()` takes a pointer to an integer variable and returns the process ID of the completed process. If no child process is running, `wait()` returns **-1**.
- Some flags that indicate the completion status of the child process are passed back with the integer pointer.

# ***How to Use wait()?***

- Wait for an unspecified child process:

```
wait(&status) ;
```

- Wait for a number, say **n**, of unspecified child processes:

```
for (i = 0; i < n; i++)  
    wait(&status) ;
```

- Wait for a specific child process whose ID is known:

```
while (pid != wait(&status))  
    ;
```

# *wait()* **System Call Example**

```
void main(void)
{
    pid_t pid, pid_child;
    int     status;

    if ((pid = fork()) == 0)    // child here
        child();
    else {                      // parent here
        parent();
        pid_child = wait(&status);
    }
}
```

# ***The exec() System Calls***

- A newly created process may run a different program rather than that of the parent.
- This is done using the `exec` system calls. We will only discuss `execvp()`:

```
int execvp(char *file, char *argv[]);
```

- ❖ `file` is a `char` array that contains the name of an executable file
- ❖ `argv[]` is the argument passed to your main program
- ❖ `argv[0]` is a pointer to a string that contains the program name
- ❖ `argv[1]`, `argv[2]`, ... are pointers to strings that contain the arguments

# *execvp() : An Example 1/2*

```
#include <stdio.h>
#include <unistd.h>

void main(void)
{
    char prog[] = { "cp" };
    char in[] = { "this.c" };
    char out[] = { "that.c" };
    char *argv[4];
    int status;
    pid_t pid;

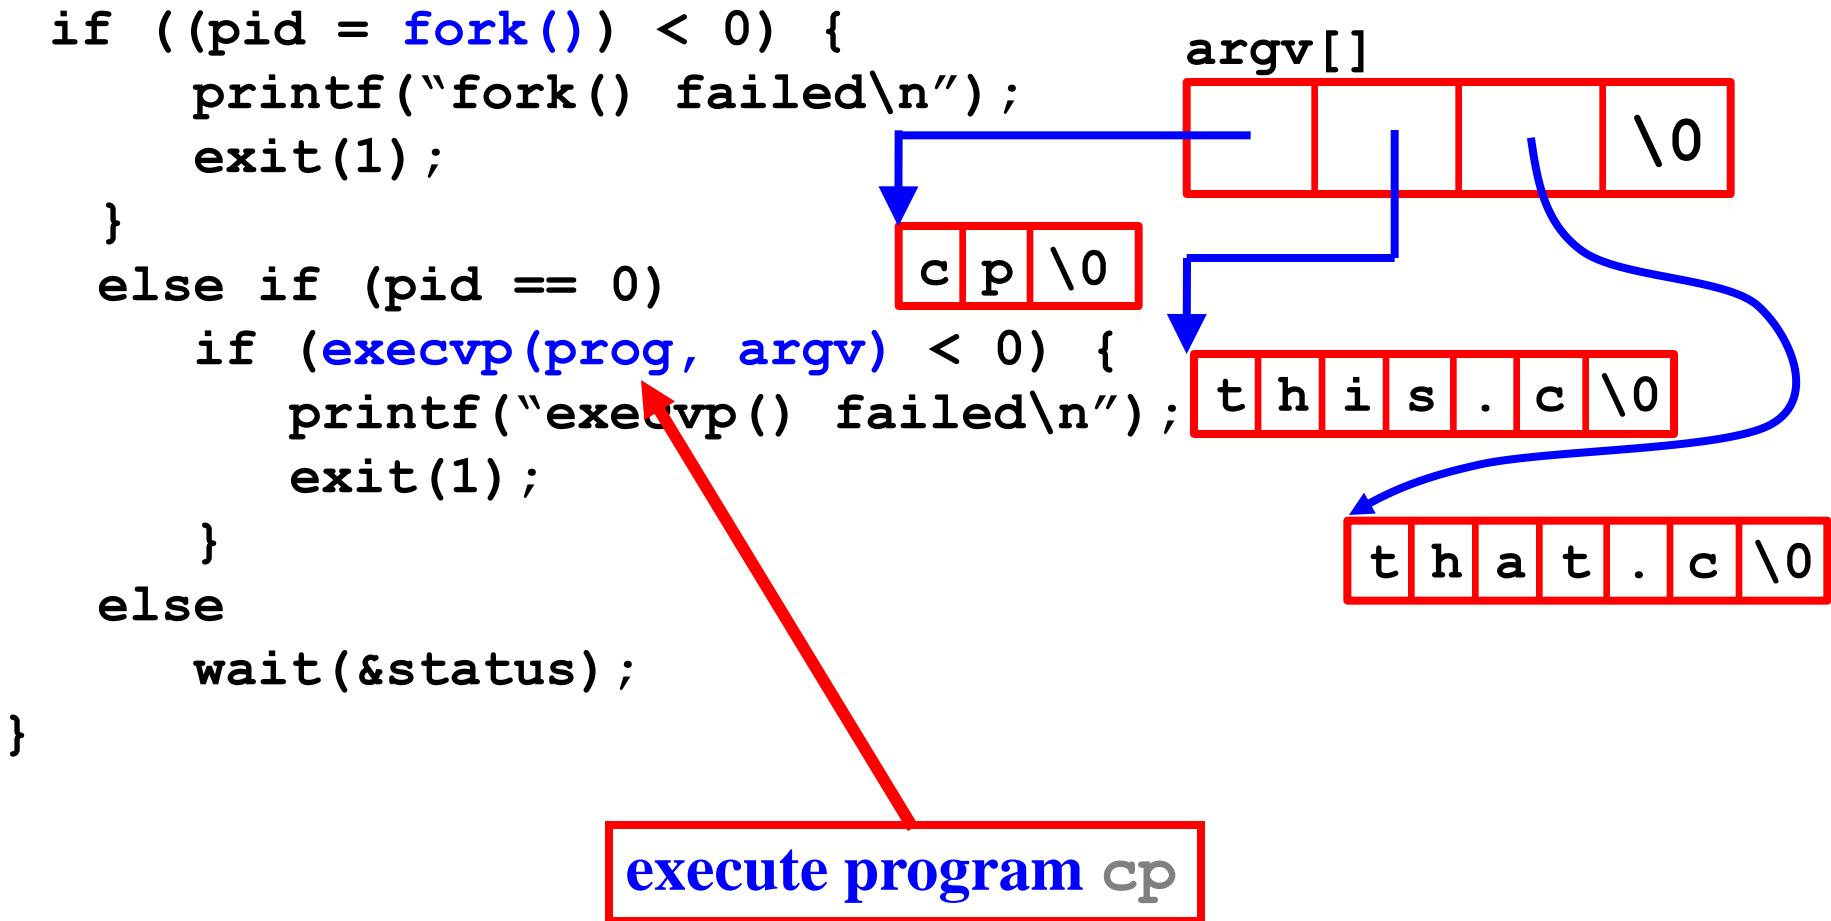
    argv[0] = prog;  argv[1] = in;
    argv[2] = out;   argv[3] = '\0';
```

The diagram illustrates the state of the `argv` array and its pointers. The `argv` array is shown as a row of four boxes, with the last box containing `\0`. Blue arrows indicate the pointers: `argv[0]` points to the first box of the `prog` array (`c p \0`), `argv[1]` points to the first box of the `in` array (`t h i s . c \0`), and `argv[2]` points to the first box of the `out` array (`t h a t . c \0`). The `in` and `out` arrays are shown as rows of boxes, with the last box containing `\0`. Green dotted arrows point from the string literals `"cp"`, `"this.c"`, and `"that.c"` in the code to their respective arrays.

*// see next slide*



# `execvp()` : An Example 2/2



# A Mini-Shell: 1/3

```
void parse(char *line, char **argv)
{
    while (*line != '\0') { // not EOLN
        while (*line == ' ' || *line == '\t' || *line == '\n')
            *line++ = '\0'; // replace white spaces with 0
        *argv++ = line; // save the argument position
        while (*line != '\0' && *line != ' '
            && *line != '\t' && *line != '\n')
            line++; // skip the argument until ...
    }
    *argv = '\0'; // mark the end of argument list
}
```

line[]

c	p		t	h	i	s	.	c		t	h	a	t	.	c	\0
---	---	--	---	---	---	---	---	---	--	---	---	---	---	---	---	----

line[]

c	p	\0	t	h	i	s	.	c	\0	t	h	a	t	.	c	\0
---	---	----	---	---	---	---	---	---	----	---	---	---	---	---	---	----

		\0
--	--	----

argv[]

# ***A Mini-Shell: 2/3***

```
void execute(char **argv)
{
    pid_t pid;
    int status;
    if ((pid = fork()) < 0) { // fork a child process
        printf("*** ERROR: forking child process failed\n");
        exit(1);
    }
    else if (pid == 0) { // for the child process:
        if (execvp(*argv, argv) < 0) { // execute the command
            printf("*** ERROR: exec failed\n");
            exit(1);
        }
    }
    else { // for the parent:
        while (wait(&status) != pid) // wait for completion
            ;
    }
}
```

# A Mini-Shell: 3/3

```
void main(void)
{
    char line[1024];    // the input line
    char *argv[64];     // the command line argument

    while (1) {         // repeat until done ....
        printf("Shell -> "); // display a prompt
        gets(line);      // read in the command line
        printf("\n");
        parse(line, argv); // parse the line
        if (strcmp(argv[0], "exit") == 0) // is it an "exit"?
            exit(0);      // exit if it is
        execute(argv);    // otherwise, execute the command
    }
}
```

**Don't forget that `gets()` is risky! Use `fgets()` instead.**

# ***What is Shared Memory?***

- The parent and child processes are run in **independent** and **separate** address spaces. All processes, parent and children included, do not share anything.
- A **shared memory segment** is a piece of memory that can be allocated and attached to an address space. Processes that have this memory segment attached will have access to it.
- But, **race conditions can occur!**

# ***Procedure for Using Shared Memory***

- Find a **key**. Unix uses this key for identifying shared memory segments.
- Use `shmget()` to allocate a shared memory.
- Use `shmat()` to attach a shared memory to an address space.
- Use `shmdt()` to detach a shared memory from an address space.
- Use `shmctl()` to deallocate a shared memory.

# Keys: 1/2

- To use shared memory, include the following:

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/shm.h>
```

- A key is a value of type `key_t`. There are three ways to generate a key:
  - ❖ Do it yourself
  - ❖ Use function `ftok()`
  - ❖ Ask the system to provide a private key.

# Keys: 2/2

- Do it yourself:

```
key_t      SomeKey;  
SomeKey = 1234;
```

- Use `ftok()` to generate one for you:

```
key_t = ftok(char *path, int ID);
```

- ❖ `path` is a path name (e.g., `"/"`)

- ❖ `ID` is an integer (e.g., `'a'`)

- ❖ Function `ftok()` returns a key of type `key_t`:

```
SomeKey = ftok("/" , 'x');
```

- Keys are **global** entities. If other processes know your key, they can access your shared memory.
- Ask the system to provide a private key with `IPC_PRIVATE`.



# Ask for a Shared Memory: 1/4

- Include the following:

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/shm.h>
```

- Use `shmget()` to request a shared memory:

```
shm_id = shmget(
    key_t key,      /* identity key */
    int size,      /* memory size */
    int flag);     /* creation or use */
```

- `shmget()` returns a shared memory ID.
- The flag, for our purpose, is either `0666` (rw) or `IPC_CREAT | 0666`. Yes, `IPC_CREAT`.

## ***Ask for a Shared Memory: 2/4***

- The following creates a shared memory of size `struct Data` with a private key `IPC_PRIVATE`. This is a creation (`IPC_CREAT`) with read and write permission (`0666`).

```
struct Data { int a; double b; char x; };  
int          ShmID;
```

```
ShmID = shmget(  
    IPC_PRIVATE,    /* private key */  
    sizeof(struct Data), /* size */  
    IPC_CREAT | 0666); /* cr & rw */
```

# ***Ask for a Shared Memory: 3/4***

- The following creates a shared memory with a key based on the current directory:

```
struct Data { int a; double b; char x;};  
int      ShmID;  
key_t    Key;  
  
Key = ftok(".", 'h');  
ShmID = shmget(  
    Key,          /* a key */  
    sizeof(struct Data),  
    IPC_CREAT | 0666);
```

# ***Ask for a Shared Memory: 4/4***

- When asking for a shared memory, the process that creates it uses `IPC_CREAT | 0666` and processes that access a created one use `0666`.
- If the return value is negative (Unix convention), the request was unsuccessful, and no shared memory is allocated.
- ***Create a shared memory before its use!***

# ***After the Execution of `shmget()`***

**Process 1**

**Process 2**

```
shmget(..., IPC_CREAT | 0666);
```



**shared memory**

**Shared memory is allocated; but, is not part of the address space**

# Attaching a Shared Memory: 1/3

- Use `shmat()` to attach an existing shared memory to an address space:

```
shm_ptr = shmat(  
    int  shm_id, /* ID from shmget() */  
    char *ptr,   /* use NULL here    */  
    int  flag); /* use 0 here        */
```

- `shm_id` is the shared memory ID returned by `shmget()`.
- Use `NULL` and `0` for the second and third arguments, respectively.
- `shmat()` returns a `void` pointer to the memory. If unsuccessful, it returns a negative integer.

# Attaching a Shared Memory: 2/3

```
struct Data { int a; double b; char x;};  
int      ShmID;  
key_t    Key;  
struct Data *p;
```

```
Key = ftok("./", 'h');
```

```
ShmID = shmget(Key, sizeof(struct Data),  
              IPC_CREAT | 0666);
```

```
p = (struct Data *) shmat(ShmID, NULL, 0);
```

```
if ((int) p < 0) {  
    printf("shmat() failed\n"); exit(1);  
}
```

```
p->a = 1; p->b = 5.0; p->x = '.';
```

# Attaching a Shared Memory: 3/3

Process 1

```
shmget(..., IPC_CREAT|0666) ;  
ptr = shmat(.....) ;
```

ptr

Process 2

```
shmget(..., 0666) ;  
ptr = shmat(.....) ;
```

ptr

shared memory

Now processes can access the shared memory



# ***Detaching/Removing Shared Memory***

- To detach a shared memory, use

**`shmdt(shm_ptr) ;`**

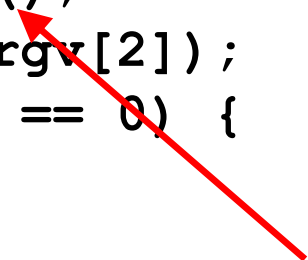
`shm_ptr` is the pointer returned by `shmat()`.

- After a shared memory is detached, it is still there. You can attach and use it again.
- To remove a shared memory, use  
**`shmctl(shm_ID, IPC_RMID, NULL) ;`**  
`shm_ID` is the shared memory ID returned by `shmget()`. After a shared memory is removed, it no longer exists.

# Communicating with a Child: 1/2

```
void main(int argc, char *argv[])
{
    int    ShmID, *ShmPTR, status;
    pid_t  pid;

    ShmID = shmget(IPC_PRIVATE, 4*sizeof(int), IPC_CREAT|0666);
    ShmPTR = (int *) shmat(ShmID, NULL, 0);
    ShmPTR[0] = getpid();          ShmPTR[1] = atoi(argv[1]);
    ShmPTR[2] = atoi(argv[2]);    ShmPTR[3] = atoi(argv[3]);
    if ((pid = fork()) == 0) {
        Child(ShmPTR);
        exit(0);
    }
    wait(&status);
    shmdt((void *) ShmPTR);  shmctl(ShmID, IPC_RMID, NULL);
    exit(0);
}
```



parent's process ID here

# Communicating with a Child: 2/2

```
void Child(int SharedMem[])
{
    printf("%d %d %d %d\n", SharedMem[0],
           SharedMem[1], SharedMem[2], SharedMem[3]);
}
```

- **Why are `shmget()` and `shmat()` not needed in the child process?**

# ***Communicating Among Separate Processes: 1/5***

- **Define the structure of a shared memory segment as follows:**

```
#define  NOT_READY  (-1)
#define  FILLED     (0)
#define  TAKEN      (1)

struct Memory {
    int  status;
    int  data[4];
};
```

# Communicating Among Separate Processes: 2/5

## The “Server”

Prepare for a shared memory



```
int main(int argc, char *argv[])  
{
```

```
    key_t          ShmKEY;  
    int            ShmID, i;  
    struct Memory  *ShmPTR;
```

```
    ShmKEY = ftok("./", 'x');  
    ShmID = shmget(ShmKEY, sizeof(struct Memory),  
                  IPC_CREAT | 0666);  
    ShmPTR = (struct Memory *) shmat(ShmID, NULL, 0);
```

# Communicating Among Separate Processes: 3/5

**shared memory not ready**  
`ShmPTR->status = NOT_READY;` **filling in data**

```
ShmPTR->data[0] = getpid();  
for (i = 1; i < 4; i++)  
    ShmPTR->data[i] = atoi(argv[i]);
```

```
ShmPTR->status = FILLED;  
while (ShmPTR->status != TAKEN)  
    sleep(1); /* sleep for 1 second */
```

```
printf("My buddy is %ld\n", ShmPTR->data[0]);
```

```
shmdt((void *) ShmPTR);  
shmctl(ShmID, IPC_RMID, NULL);  
exit(0);
```

```
}
```

**wait until the data is taken**

**detach and remove shared memory**

# Communicating Among Separate Processes: 4/5

```
int main(void)
{
```

The “Client”

```
    key_t          ShmKEY;
    int             ShmID;
    struct Memory   *ShmPTR;
```

prepare for shared memory



```
    ShmKEY=ftok("./", 'x');
```

```
    ShmID = shmget(ShmKEY, sizeof(struct Memory), 0666);
```

```
    ShmPTR = (struct Memory *) shmat(ShmID, NULL, 0);
```

```
    while (ShmPTR->status != FILLED)
```

```
        ;
```

```
    printf("%d %d %d %d\n", ShmPTR->data[0],
```

```
        ShmPTR->data[1], ShmPTR->data[2], ShmPTR->data[3]);
```

```
    ShmPTR->data[0] = getpid();
```

```
    ShmPTR->status = TAKEN;
```

```
    shmdt((void *) ShmPTR);
```

```
    exit(0);
```

```
}
```

# ***Communicating Among Separate Processes: 5/5***

- The “server” must run first to **prepare** a shared memory.
- Run the server in one window, and run the client in another a little later.
- Or, run the server as a background process. Then, run the client in the foreground later:

```
server 1 3 5 &  
client
```

- This version uses **busy waiting**.
- **One may use Unix semaphores for mutual exclusion.**

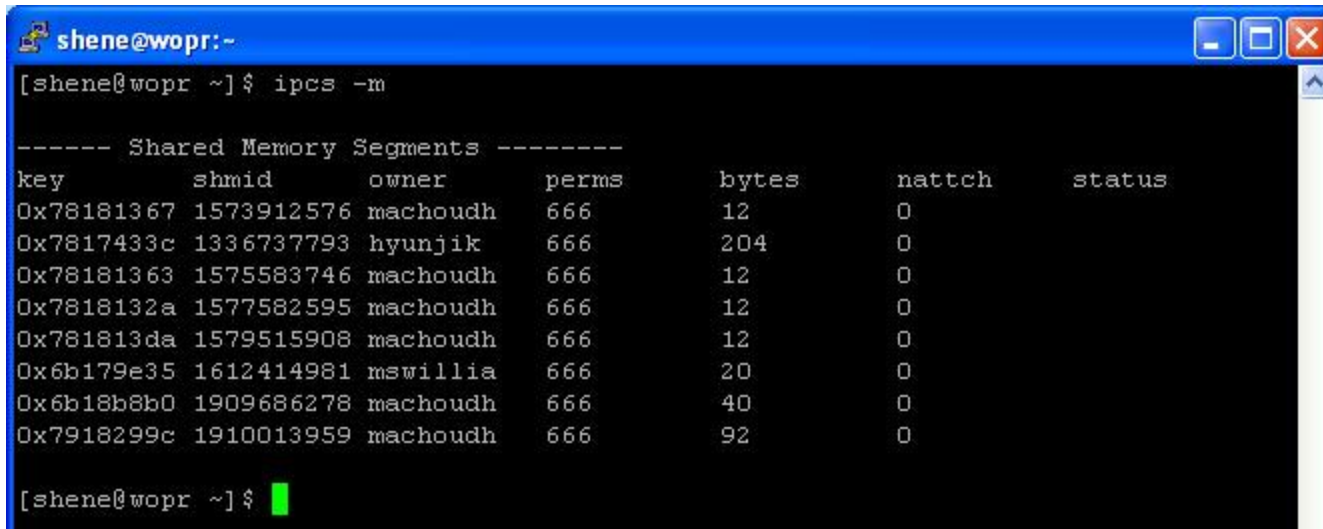


## ***Important Notes: 1/3***

- If you do not remove your shared memory segments (e.g., program crashes before the execution of `shmctl()`), they will be in the system forever until the system is shut down. This will degrade the system performance.
- Use the `ipcs` command to check if you have shared memory segments left in the system.
- Use the `ipcrm` command to remove your shared memory segments.

## Important Notes: 2/3

- To see existing shared memory segments in the system, use `ipcs -m`, where `m` means shared memory.
- The following is a snapshot on `wopr`:



```
shene@wopr:~$ ipcs -m
```

key	shmid	owner	perms	bytes	nattch	status
0x78181367	1573912576	machoudh	666	12	0	
0x7817433c	1336737793	hyunjik	666	204	0	
0x78181363	1575583746	machoudh	666	12	0	
0x7818132a	1577582595	machoudh	666	12	0	
0x781813da	1579515908	machoudh	666	12	0	
0x6b179e35	1612414981	mswillia	666	20	0	
0x6b18b8b0	1909686278	machoudh	666	40	0	
0x7918299c	1910013959	machoudh	666	92	0	

```
[shene@wopr ~]$
```

## ***Important Notes: 3/3***

- To remove a shared memory, use the `ipcrm` command as follows:
  - ❖ `ipcrm -M shm-key`
  - ❖ `ipcrm -m shm-ID`
- You have to be the owner (or super user) to remove a shared memory.

**The End**