

CS 3411 Systems Programming

Department of Computer Science
Michigan Technological University

File Systems and Protection


Today's Topics

- ▶ Removable File Systems
- ▶ Special Files
- ▶ Links (Hard and Symbolic)
- ▶ FIFOs
- ▶ Protection
- ▶ Opening files via kernel calls!


Removable File Systems

- ▶ Only the `/` directory is known to Unix at boot time
- ▶ Any directory is capable of having a file system mounted on it
- ▶ Usually restricted to privileged accounts

`/dev/sda1`



`/dev/sda2`



Special Files

- ▶ Makes physical devices appear to be part of file system hierarchy
- ▶ Provides uniform I/O interface (Can read or write just like an ordinary file!)
- ▶ Read/Write maps onto direct I/O on the device
- ▶ `/dev/sda` **or** `/dev/sda0` **or** `/dev/hdb2` **or** ...
- ▶ `/dev/tty`, `/dev/mem`, `/dev/kmem`, `/dev/null`

Other File Types

Links

- ▶ Either allows one file to point to another file (Symbolic Links)
- ▶ Or file is accessible via multiple names (Hard Links)

Sockets

- ▶ Sockets are interfaces to a network

FIFOs (Named Pipes)

- ▶ FIFOs are mechanisms for communication between unrelated processes

Symbolic Links Example

```
$ echo "This_is_a_test_file." > test_sl.txt
$ cat test_sl.txt
This is a test file.
$ ln -s test_sl.txt test_link
$ ls -la test_sl.txt
-rw-r--r-- 1 galolu cs-stu 21 May 16 15:11 test_sl.txt
$ ls -la test_link
lrwxrwxrwx 1 galolu cs-stu 11 May 16 15:12 test_link -> test_sl.txt
$ cat test_link
This is a test file.
$ echo "Appending_data." >> test_link
$ cat test_link
This is a test file.
Appending data.
$ rm test_sl.txt
$ cat test_link
cat: test_link: No such file or directory
```

FIFO Example

```
$ mkfifo my_fifo
$ ls -l my_fifo
prw-r--r-- 1 galolu cs-stu 0 May 16 15:21 my_fifo
$ echo "Test_String" > test_file
$ wc < my_fifo &
[1] 5972
$ cat test_file | tee my_fifo | grep e
Test String
1      2      12
```

Protection

- ▶ User and user's processes have associated *user id* and *group id*
- ▶ The Unix file system has 12 bits to specify protection
- ▶ These bits are specified by the creator of the file
- ▶ `_ _ _ r w x r w x r w x`
 - ▶ First three bits discussed later
 - ▶ Second group of 3 bits apply if the user is owner of the file
 - ▶ Third group of 3 bits apply if the user is not the owner, but is in the group of the owner
 - ▶ Last group of 3 bits apply if user is not owner, and user is not in group of owner
- ▶ For ordinary files: `r` is read, `w` is write and `x` is execute
- ▶ For directories: `r` allows `ls`, `w` allows file creation and deletion and `x` allows pass-through
- ▶ `chmod` allows the owner to change protection of file
- ▶ Protection bits are expressed as 4 *octal* digits

Protection

- ▶ Processes have *real* and *effective* user ids
- ▶ Effective user id is used in protection validation
- ▶ The leftmost protection bit is for `suid` (set user id)
- ▶ If a binary with the `suid` bit set is executed, it will execute with the privilege level of its owner, not the invoker
- ▶ The second bit is `sgid` (set group id), and does the same thing as `suid` for groups instead of users
- ▶ The third bit is the sticky bit.
 - ▶ If on an executable, it will tell the OS to keep the "text" segment of the file in swap so it can be re-executed rapidly
 - ▶ If on a directory, it means people who have write access to the directory can create files, but no one can delete files but the owner

Protection Example

```
$ mkdir newdir
$ ls -l
total 168
-rw-r--r-- 1 galolu cs-stu 4754 May 16 14:59 mount.pdf
prw-r--r-- 1 galolu cs-stu    0 May 16 15:21 my_fifo
drwxr-xr-x 2 galolu cs-stu  80 May 16 15:42 newdir
-rw-r--r-- 1 galolu cs-stu 2089 May 16 15:41 slides05.aux
...
$ umask
0022
$ echo "Test" > newdir/newFile
$ cat newdir/newFile
Test
$ ls newdir
newFile
$ chmod u-r newdir
$ ls newdir
ls: cannot open directory newdir: Permission denied
$ cat newdir/newFile
Test
$ chmod u-x newdir
$ cat newdir/newFile
cat: newdir/newFile: Permission denied
```

open() Kernel Call

- ▶ Read the manual!
- ▶ We have to tell the function what file we want opened, and with what intent.
- ▶ Some of the possible flags are:
 - ▶ `O_RDONLY` - Open for reading only
 - ▶ `O_WRONLY` - Open for writing only
 - ▶ `O_RDWR` - Open for reading and writing
 - ▶ `O_APPEND` - If set, the seek pointer will be set to the end of the file prior to each write.
 - ▶ `O_CREAT` - If the file exists, this flag has no effect. Otherwise, the file is created, owner ID of the file set to the *effective* user ID of the process
 - ▶ `O_TRUNC` - If the file exists and is a regular file, and the file is successfully opened via `O_RDWR` or `O_WRONLY`, its length is truncated to zero but the mode and owner are unchanged.

open() Kernel Call

- ▶ A successful open call returns a *file descriptor*
- ▶ File descriptors, when used to do I/O, access a *currency indicator* stored by the OS to figure out where to start the next I/O operation
- ▶ If you open the same file multiple times, you will end up with multiple file descriptors and currency indicators!
- ▶ Some sample calls:
- ▶ `datafile = open("mydata", O_RDWR);`

```
newfile = open("mydata", O_RDWR|O_CREAT, 0644);
```

open() vs. fopen()

- ▶ `open()` is a Unix kernel call; `fopen()` is a standard I/O library function
- ▶ `open()` returns an `int` (file descriptor); `fopen()` returns a `FILE *` (an stdio stream)
- ▶ Unix I/O kernel calls [`read()`, `write()`, `lseek()`] operate on file descriptors; standard I/O functions [`fscanf()`, `fprintf()`, `fread()`, `fwrite()`, `fseek()`] operate on stdio streams
- ▶ User level application programs should use `fopen()` for conventional files; systems programs should use `open()`.
- ▶ Never mix and match the two calls!

Next Week

- ▶ First assignment!
- ▶ Other kernel call I/O functions
- ▶ File locking
- ▶ Processes