

# CS 3411 Systems Programming

Department of Computer Science  
Michigan Technological University

## File Systems (cont.)

# Today's Topics

- ▶ File manipulation via kernel calls
- ▶ File locking and synchronous writing

# open() Kernel Call

- ▶ Read the manual!
- ▶ We have to tell the function what file we want opened, and with what intent.
- ▶ Some of the possible flags are:
  - ▶ `O_RDONLY` - Open for reading only
  - ▶ `O_WRONLY` - Open for writing only
  - ▶ `O_RDWR` - Open for reading and writing
  - ▶ `O_APPEND` - If set, the seek pointer will be set to the end of the file prior to each write.
  - ▶ `O_CREAT` - If the file exists, this flag has no effect. Otherwise, the file is created, owner ID of the file set to the *effective* user ID of the process
  - ▶ `O_TRUNC` - If the file exists and is a regular file, and the file is successfully opened via `O_RDWR` or `O_WRONLY`, its length is truncated to zero but the mode and owner are unchanged.

# open() Kernel Call

- ▶ A successful open call returns a *file descriptor*
- ▶ File descriptors, when used to do I/O, access a *currency indicator* stored by the OS to figure out where to start the next I/O operation
- ▶ If you open the same file multiple times, you will end up with multiple file descriptors and currency indicators!
- ▶ Some sample calls:
- ▶ `datafile = open("mydata", O_RDWR);`

```
newfile = open("mydata", O_RDWR|O_CREAT, 0644);
```

# open() vs. fopen()

- ▶ `open()` is a Unix kernel call; `fopen()` is a standard I/O library function
- ▶ `open()` returns an `int` (file descriptor); `fopen()` returns a `FILE *` (an stdio stream)
- ▶ Unix I/O kernel calls [`read()`, `write()`, `lseek()`] operate on file descriptors; standard I/O functions [`fscanf()`, `fprintf()`, `fread()`, `fwrite()`, `fseek()`] operate on stdio streams
- ▶ User level application programs should use `fopen()` for conventional files; systems programs should use `open()`.
- ▶ Never mix and match the two calls!

# Default File Descriptors

- ▶ 0 is stdin
- ▶ 1 is stdout
- ▶ 2 is stderr
- ▶ You can use these as integer values directly without opening them
- ▶ Opened on process creation

# read() Kernel Call

- ▶ Manual page!
- ▶ `ssize_t read(int fd, void *buf, size_t count);`
- ▶ Attempts to read up to count bytes from file descriptor fd into the buffer starting at buf
- ▶ Normal use:

```
int datafile;  
ssize_t num;  
char buff[100];  
datafile = open("mydata", O_RDWR);  
num = read(datafile, buff, 100);
```

# read() Kernel Call

- ▶ Use sizeof when reading into variables
- ▶ Check the output to confirm read was done as you wanted:
  - ▶ 0 means end of file
  - ▶ -1 means an error
  - ▶ Anything else is the number of bytes read
- ▶ We could also read() to read saved structs from disk



# write() Kernel Call

- ▶ Manual page!
- ▶ `ssize_t write(int fd, const void *buf, size_t count);`
- ▶ Writes up to count bytes to the file referenced by the file descriptor from the buffer starting at buf
- ▶ Normal use:

```
int datafile;  
ssize_t num;  
char buff[100];  
buff = "Some_string ...";  
datafile = open("mydata", O_RDWR);  
num = write(datafile, buff, 100);
```

## write() Kernel Call

- ▶ Use sizeof when writing variables to files
- ▶ Check the output to confirm read was done as you wanted:
  - ▶ 0 means nothing was written
  - ▶ -1 means an error
  - ▶ Anything else is the number of bytes written
- ▶ We could also write() to write structs to disk

# lseek() Kernel Call

- ▶ Manual page!
- ▶ `off_t lseek(int fd, off_t offset, int whence);`
- ▶ Repositions the currency indicator of the open file fd to the argument offset according to whence
- ▶ If whence is SEEK\_SET, the currency indicator is set to offset bytes
- ▶ If whence is SEEK\_CUR, the currency indicator is set to current location plus offset bytes
- ▶ If whence is SEEK\_END, the currency indicator is set to the size of the file plus offset bytes

# lseek() Kernel Call

- Moves currency indicator without doing any I/O

```
n = lseek(fd, (off_t)100, SEEK_SET);  
n = lseek(fd, (off_t)100, SEEK_CUR);  
n = lseek(fd, (off_t)-100, SEEK_END);  
n = lseek(fd, (off_t)100, SEEK_END);
```

# File Locking

- ▶ Each process has it's own pointer to the current position and file status flags, possibly to the same file
- ▶ `read()` and `write()` are atomic from perspective of other processes
- ▶ When file is opened for append, position pointer moved to end of file prior to `write()`
- ▶ While `read()` and `write()` are atomic, concurrent operations on shared files can have unexpected results:

# Examples

```
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>
#include <string.h>

main() {
    char buf[] = "aaa_aaa_aaa_aaa_aaa_aaa_aaa_aaa_aaa\n";
    int sharedFd, i, j, x;

    sharedFd = open("test.out", O_RDWR|O_APPEND|O_CREAT, 0644);
    for (i=0; i < strlen(buf); i++) {
        write(sharedFd, &buf[i], 1);
        for (j = 0; j < 100000; j++) {
            x = 0;
        }
    }
}
```

# Examples

```
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>
#include <string.h>

main() {
    char buf[] = "bbb_bbb_bbb_bbb_bbb_bbb_bbb_bbb_bbb\n";
    int sharedFd, i, j, x;

    sharedFd = open("test.out", O_RDWR|O_APPEND|O_CREAT, 0644);
    for (i=0; i < strlen(buf); i++) {
        write(sharedFd, &buf[i], 1);
    }
}
```

# Examples

Append is not the only problem!

```
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>
#include <string.h>

main() {
    char buf[] = "aaa_aaa_aaa_aaa_aaa_aaa_aaa_aaa\n";
    int sharedFd, i, j, x;

    sharedFd = open("test.out", O_RDWR|O_CREAT, 0644);
    srandom(getpid());
    for (i=0; i < strlen(buf); i++) {
        x = random() % 3;
        sleep(x);
        write(sharedFd, &buf[i], 1);
    }
}
```



# Examples

```
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>
#include <string.h>

main() {
    char buf[] = "bbb_bbb_bbb_bbb_bbb_bbb_bbb_bbb_bbb\n";
    int sharedFd, i, j, x;

    sharedFd = open("test.out", O_RDWR|O_CREAT, 0644);
    srandom(getpid());
    for (i=0; i < strlen(buf); i++) {
        x = random() % 3;
        sleep(x);
        write(sharedFd, &buf[i], 1);
    }
}
```

# Locking Portions of Files

- ▶ `int fcntl(int fd, int cmd, ... /* arg */);`
- ▶ `fcntl()` performs an operation described by the `cmd` argument on the open file descriptor `fd`
- ▶ `cmd` options for locking
  - ▶ `F_GETLK` returns information about the lock if lock is held, otherwise returns `F_UNLCK` in the appropriate field of the struct (next slide)
  - ▶ This does not acquire the lock!
  - ▶ `F_SETLK` sets the lock as described
  - ▶ `F_SETLKW` tries to set the lock as described, but will put the calling process to sleep if lock can't be granted

# The flock Struct

```
struct flock {  
    short l_type; /* Type of lock: F_RDLCK, F_WRLCK, F_UNLCK */  
    short l_whence; /* How to interpret l_start: SEEK_SET,  
    SEEK_CUR, SEEK_END */  
    off_t l_start; /* Starting offset for lock */  
    off_t l_len; /* Number of bytes to lock */  
    pid_t l_pid; /* PID of process blocking our lock,  
    returned with F_GETLCK */  
}
```

These are advisory locks! They will only work if all processes honor lock status.

# Examples I

```
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>
#include <string.h>

main() {
    char buf[] = "aaa_aaa_aaa_aaa_aaa_aaa_aaa_aaa\n";
    int sharedFd, i, j, x;
    struct flock sharedFdLock;

    sharedFd = open("test.out", O_RDWR|O_APPEND|O_CREAT, 0644);
    sharedFdLock.l_type=F_WRLCK;
    sharedFdLock.l_start=0;
    sharedFdLock.l_whence=SEEK_SET;
    sharedFdLock.l_len=0;

    fcntl(sharedFd, F_SETLKW, &sharedFdLock);
    write(2, "Writer_1_beginning_write.\n", 26);
    for (i=0; i < strlen(buf); i++) {
        write(sharedFd, &buf[i], 1);
        for (j = 0; j < 100000; j++) {
            x = 0;
        }
    }
}
```

# Examples II

```
sharedFdLock.l_type=F_UNLCK;  
sharedFdLock.l_start=0;  
sharedFdLock.l_whence=SEEK_SET;  
sharedFdLock.l_len=0;  
  
write(2, "Writer_1_done_writing.\n", 23);  
fcntl(sharedFd,F_SETLK, &sharedFdLock);  
sleep(2);  
}
```

# Examples I

```
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>
#include <string.h>

main() {
    char buf[] = "bbb_bbb_bbb_bbb_bbb_bbb_bbb_bbb\n";
    int sharedFd, i, j, x;
    struct flock sharedFdLock;

    sharedFd = open("test.out", O_RDWR|O_APPEND|O_CREAT, 0644);
    sharedFdLock.l_type=F_WRLCK;
    sharedFdLock.l_start=0;
    sharedFdLock.l_whence=SEEK_SET;
    sharedFdLock.l_len=0;

    fcntl(sharedFd, F_SETLKW, &sharedFdLock);
    write(2, "Writer_2_beginning_write.\n", 26);

    for (i=0; i < strlen(buf); i++) {
        write(sharedFd, &buf[i], 1);
    }

    sharedFdLock.l_type=F_UNLCK;
```

# Examples II

```
sharedFdLock.l_start=0;
sharedFdLock.l_whence=SEEK_SET;
sharedFdLock.l_len=0;

write(2, "Writer_2_done_writing.\n", 23);
fcntl(sharedFd,F_SETLK, &sharedFdLock);
sleep(2);
}
```