

CS 3411 Systems Programming

Department of Computer Science
Michigan Technological University

Unix Processes (cont.)

Today's Topics

- ▶ More about `execve`
- ▶ Debugging

fork()

- ▶ A forked child inherits open files of the parent
- ▶ The child process descriptor is a copy of the parent's process descriptor, except:
 - ▶ Return value from `fork()`
 - ▶ PID, PPID
 - ▶ Pending signals and alarms
 - ▶ File locks
 - ▶ Execution times

Executing a New Binary

- ▶ `execve()` is used to execute a new program
- ▶ Manual page!
- ▶ This function executes the program it is pointed to
- ▶ On success, `execve()` does not return: The process calling `execve()` is completely replaced by the newly executed process
- ▶ On error, -1 is returned
- ▶ File descriptors may be set to close on exec!

Creating a New Process

- ▶ Exec is most useful when used with fork
- ▶ In Unix, a new process is created by first forking an existing process, then calling a variant of exec from there
- ▶ Most process attributes are preserved, including the PID, PPID, file locks, pending signals, execution times and open files

execve() Example

```
#include <stdio.h>
#include <errno.h>
#include <stdlib.h>

main () {
    char *a[4], *e[3];

    a[0] = "child";
    a[1] = "argument1";
    a[2] = "argument2";
    a[3] = NULL;

    e[0] = "ENV0=val0";
    e[1] = "ENV1=val1";
    e[2] = NULL;

    execve("child", a, e);
    /* If we get here, something went wrong */
    perror("parent1");
    exit(1);
}
```

execve() Example

```
#include <stdio.h>
```

```
main(argc, argv, envp)
```

```
int argc;
```

```
char *argv[], *envp[];
```

```
{
```

```
    int i;
```

```
    char **ep;
```

```
    printf("child is running\n");
```

```
    for (i = 0; i < argc; i++) {
```

```
        printf("argv[%d]=%s\n", i, argv[i]);
```

```
    }
```

```
    for (ep = envp; *ep; ep++) {
```

```
        printf("%s\n", *ep);
```

```
    }
```

```
}
```

More Examples

```
#include <fcntl.h>
#include <stdio.h>
#include <unistd.h>
#include <stdlib.h>
main(argc, argv)
int argc; char *argv[];
{
    int forkid, charnum;
    char fdval[20];
    if (argc != 3) {
        fprintf(stderr, "Usage: _pexec_ filename _charnum\n");
        exit(1);
    }
    if ((forkid = open(argv[1], O_RDONLY)) < 0) {
        fprintf(stderr, "Cannot _open_%s\n", argv[1]);
        exit(2);
    }
    sprintf(fdval, "%d", forkid); /* sprintf! */
    if (fork() == 0) {
        execl("pchild", "pchild", fdval, argv[2], (char *)0);
        fprintf(stderr, "Unable _to _exec\n");
        exit(3);
    }
    printf("Parent _is _after _fork/exec\n");
}
```


More Examples

```
#include <sys/types.h>
#include <unistd.h>
#include <stdio.h>
#include <stdlib.h>

main(argc, argv)
int argc; char *argv[];
{
    int myfd;
    char gotch, val;
    if (argc != 3) {
        fprintf(stderr, "Usage: _pchild _filename _charnum\n");
        exit(1);
    }
    myfd = atoi(argv[1]);
    gotch = atoi(argv[2]);
    lseek(myfd, (off_t)gotch, SEEK_SET);
    read(myfd, &val, 1);
    printf("Child _got _char_%d _from _fd _%d: _%c\n", gotch, myfd, val);
}
```

More Examples

```
#include <fcntl.h>
#include <stdio.h>
#include <unistd.h>
#include <stdlib.h>
main(argc, argv)
int argc; char *argv[];
{
    int forkid, charnum;
    if (argc != 3) {
        fprintf(stderr, "Usage: _pioexec _filename _charnum\n");
        exit(1);
    }
    if ((forkid = open(argv[1], O_RDONLY)) < 0 ) {
        fprintf(stderr, "Cannot _open _%s\n", argv[1]);
        exit(2);
    }
    if (fork() == 0) {
        close(0); dup(forkid); close(forkid);
        execl("piochild", "piochild", argv[2], (char *)0);
        fprintf(stderr, "Unable _to _exec\n");
        exit(3);
    }
    printf("Parent _is _after _fork/exec\n");
}
```

More Examples

```
#include <sys/types.h>
#include <unistd.h>
#include <stdio.h>
#include <stdlib.h>

main(argc, argv)
int argc; char *argv[];
{
    int myfd;
    char gotch, val;
    if (argc != 2) {
        fprintf(stderr, "Usage: _piokid _charnum\n");
        exit(1);
    }
    gotch = atoi(argv[1]);
    lseek(0, (off_t)gotch, SEEK_SET);
    read(0, &val, 1);
    printf("Child _got _char _%d _from _stdin: _%c\n", gotch, val);
}
```

File Descriptor Example

```
#include <sys/types.h>
#include <unistd.h>
#include <sys/stat.h>
#include <fcntl.h>

main() {
    pid_t cpid;
    int x, fd;
    char ch;

    x = 5;
    cpid = fork();
    if (cpid == 0) {
        fd = open("aFile", O_RDWR|O_CREAT, 0644);
        x++; ch=x+48;
        write(fd,&ch,1);
    } else {
        fd = open("aFile", O_RDWR|O_CREAT, 0644);
        x++; ch=x+48;
        write(fd,&ch,1);
    }
}
```

More Examples

How would we need to code to get the following process structure:

- ▶ Parent
 - ▶ Child 1
 - ▶ Child 2

More Examples

How would we need to code to get the following process structure:

- ▶ Parent
 - ▶ Child 1
 - ▶ Child 2

More Examples

How would we need to code to get the following process structure:

- ▶ Parent
 - ▶ Child 1
 - ▶ Child 2

Exec Example

```
#include <sys/types.h>
#include <unistd.h>

main() {
    pid_t cpid;
    int i;

    for(i = 0; i < 2; i++) {
        cpid=fork();
        if (cpid==0) execl("bogus", "bogus", (char *)0);
    }
}
```


Exec Example

```
#include <sys/types.h>
#include <unistd.h>

main() {
    pid_t cpid;
    int i;

    for(i = 0; i < 2; i++) {
        cpid=fork();
        if (cpid==0) {
            execl("bogus", "bogus", (char *)0);
            exit(1);
        }
    }
}
```

Parent/Child Synchronization (wait/exit)

- ▶ The `exit()` and `_exit()` calls
- ▶ As usual, take a look at the manual page.
- ▶ Terminates the calling process immediately
- ▶ As a convention: exit status of 0 is *normal* termination
- ▶ Any other status denotes an error or an exceptional condition on termination

wait() call

- ▶ The wait() call
- ▶ The manual!
- ▶ A parent process is obligated to wait for its children to exit
- ▶ Suspends execution of the current process until a child has exited

No Waiting

```
#include <unistd.h>
#include <stdlib.h>
main() {
    if (fork() == 0) {
        exit(1);
    } else {
        sleep(20);
        exit(1);
    }
}
```

With Waiting

```
#include <unistd.h>
#include <stdlib.h>
#include <stdio.h>
#include <sys/types.h>
#include <sys/wait.h>
main() {
    int status;
    if (fork() == 0) {
        sleep(20);
        exit(51);
    } else {
        printf("pid=_%d\n", wait(&status));
        printf("status=_%x\n", status);
        if (WIFEXITED(status))
            printf("Status_(via_macro):_%d\n", WEXITSTATUS(status));
        exit(0);
    }
}
```

Debugging

- ▶ We'll focus on gcc and gdb
- ▶ We'll also take a look at ddd, which is a GUI for gdb and various other debuggers
- ▶ What is debugging?
 - ▶ Program *state* is a snapshot of all variables, PC, etc.
 - ▶ A statement in your program transforms one program state into another
 - ▶ You should be able (at some level) to express what you expect the state of your program to be after every statement
 - ▶ Often state predicates on program state; i.e., “if control is here, I expect the following to be true.”
- ▶ Let's look at a toy example

Sample Program

```
#include <stdio.h>
int sum=0, val, num=0;
double ave;
/* sum and num should be 0 */
main()
{
    /* sum should be the total of the num
    input values processed */
    while (scanf("%d\n",&val) != EOF) {
        sum += val;
        num++;
    }
    /* sum should be the total of the num
    input values and there is no more input */
    if (num > 0) {
        ave= sum/num;
        /* ave should be the floating point mean of
        the num input data values */
        printf("Average is %f\n", ave);
    }
}
```

Using gdb

- ▶ Make sure to compile source with the `-g` switch asserted.
- ▶ In our case, `gcc -g ave.c`
- ▶ Breakpoint: line in source code at which debugger will pause execution.
- ▶ At breakpoint, can examine values of relevant components of program state. `break` command sets a breakpoint; `clear` removes the breakpoint.
- ▶ Diagnostic `printf()` crude way of getting a snapshot of program state at a given point.
- ▶ Once paused at a breakpoint, use `gdb print`, or `display` to show variable or expression values. `display` will automatically print values when execution halts at breakpoint.
- ▶ From a breakpoint, may `step` or `next` to single step the program. `step` stops after next source line is executed. `next` similar, but executes functions without stopping.

Using gdb

- ▶ Can find out where execution is, in terms of function calls, with the `where` command.
- ▶ Let's play with the toy program!
- ▶ We'll put the buggy set in a data file.
- ▶ We can also debug post mortem in crashing programs (`bintree.c`)!
- ▶ May need to enable `ulimit -c`

A GUI for gdb: ddd

- ▶ The ddd program is a front-end for gdb
- ▶ Can use the mouse to set breakpoints!
- ▶ Values are graphically displayed
- ▶ Can visualize complex structures

Debugging Tips

- ▶ Examine the most recent change
- ▶ Debug it now, not later
- ▶ Read before typing
- ▶ Make the bug reproducible
- ▶ Display output to localize your search
- ▶ Write a log file
- ▶ Use tools
- ▶ Keep records