

CS 3411 Systems Programming

Department of Computer Science
Michigan Technological University

Pipe Inter-Process Communication in Unix

Today's Topics

- ▶ How to communicate between processes without using signals
- ▶ Creating and Using Pipes

How to Communicate Between Processes

- ▶ Have two primitive mechanisms in hand:
 - ▶ exit/wait
 - ▶ signals
- ▶ The sample program we'll be working on:
 - ▶ Parent creates child
 - ▶ Send child an int, x
 - ▶ Child computes $2x$ and returns result
 - ▶ Maybe a loop that keeps doing this!
- ▶ Parent and child share no memory; i.e. no common variables through which to communicate

How to Communicate Between Processes

- ▶ Parent and child share open file descriptors, even after `fork()` or `fork()/exec()`
- ▶ Proposed solution:
 - ▶ Parent opens files, forks child
 - ▶ Child arranges descriptors with `close` and `dup`, execs new binary
 - ▶ Parent and child communicate via `read()/write()` on shared file

Example Code

```
#include <fcntl.h>
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>
#include <stdlib.h>
main() {
    int fd, val, dblval;
    fd = open("commofile", O_RDWR|O_CREAT|O_TRUNC, 0644);
    if (fork() == 0) { /* CHILD */
        read(fd, &val, sizeof(int));
        lseek(fd, 0, SEEK_SET);
        dblval = 2 * val;
        write(fd, &dblval, sizeof(int));
        lseek(fd, 0, SEEK_SET);
        exit(0);
    } else { /* PARENT */
        val = 2;
        fprintf(stderr, "Asking child to double %d\n", val);
        write(fd, &val, sizeof(int));
        lseek(fd, 0, SEEK_SET);
        read(fd, &dblval, sizeof(int));
        fprintf(stderr, "Child replied with %d\n", dblval);
        wait(NULL);
        exit(0);
    }
}
```

Results

- ▶ Why did it turn out like that?
- ▶ Fundamental problem: Need more control over access to the shared file
- ▶ Specifically, read from an empty file (or read when currency indicator is at EOF) should delay caller until data is available to read
- ▶ The Unix solution is a construct called a *pipe*

The pipe() System Call

- ▶ The `pipe()` function takes one argument in the form of an array with 2 elements
- ▶ The first element is the *read* end of the pipe
- ▶ The second element is the *write* end of the pipe
- ▶ Data written to the pipe is buffered by the kernel until it is read

Pipes

- ▶ Can think of a pipe as an unnamed, fixed length file maintained by the kernel
- ▶ We also looked at the named versions! They're identical once opened - no data is to the device in either case!
- ▶ Pipes have separate file descriptors as well as currency indicators for reading and writing
- ▶ Some special properties:
 - ▶ A read from a pipe that doesn't have sufficient data to satisfy the read will block the reader until the data is available
 - ▶ A write to a pipe that is full will delay the writer until space becomes available
 - ▶ From a reader's perspective, EOF can only happen if there is no data in the pipe and all write descriptors on the pipe are closed
 - ▶ From a writer's perspective, an attempt to write to a pipe without live read descriptors will result in a SIGPIPE signal being sent to the writer

Example Code

```
#include <fcntl.h>
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>
#include <stdlib.h>
int fd[2], val = 0, dblval = 0;
main() {
    pipe(fd);
    if (fork() == 0) { /* CHILD */
        while(read(fd[0], &val, sizeof(int)) != 0) {
            dblval = 2*val;
            write(fd[1], &dblval, sizeof(int));
        }
        exit(0);
    } else { /* PARENT */
        for (val = 1; val <= 3; val++) {
            fprintf(stderr, "Asking child to double %d\n", val);
            write(fd[1], &val, sizeof(int));
            read(fd[0], &dblval, sizeof(int));
            fprintf(stderr, "Child replied with %d\n", dblval);
        }
        wait(NULL);
    }
}
```

Results

- ▶ What now?
- ▶ Another problem: Pipe is essentially a construct for *unidirectional* transfer of information. Parent is reading its own data written into the pipe.
- ▶ Need one pipe for synchronized parent-to-child communication
- ▶ And a second pipe for synchronized child-to-parent communication
- ▶ When do we get EOF?

Example Code

```
#include <fcntl.h>
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>
#include <stdlib.h>
int p2c[2], c2p[2], val = 0, dblval = 0;
main() {
    pipe(p2c); pipe(c2p);
    if (fork() == 0) { /* CHILD */
        close(p2c[1]); close(c2p[0]);
        while(read(p2c[0], &val, sizeof(int)) != 0) {
            dblval = 2*val;
            write(c2p[1], &dblval, sizeof(int));
        }
        exit(0);
    } else { /* PARENT */
        close(c2p[1]); close(p2c[0]);
        for (val = 1; val <= 3; val++) {
            fprintf(stderr, "Asking child to double %d\n", val);
            write(p2c[1], &val, sizeof(int));
            read(c2p[0], &dblval, sizeof(int));
            fprintf(stderr, "Child replied with %d\n", dblval);
        }
        close(p2c[1]); close(c2p[0]);
        wait(NULL);
    }
}
```

Results

- ▶ It works as expected this time!
- ▶ Let's take a look at some other examples

Real Examples I

```
#include <fcntl.h>
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>
#include <stdlib.h>
/* We want to execute "ps -aux | grep root" */
main() {
    int isParent;
    int apipe[2];
    char *cmd[2][3];

    cmd[0][0] = "ps";    cmd[0][1] = "-aux"; cmd[0][2] = NULL;
    cmd[1][0] = "grep";  cmd[1][1] = "root"; cmd[1][2] = NULL;

    pipe(apipe);
    isParent = fork();

    if (!isParent) { /* Child is going to be "grep root" */
        /* We want stdin connected to our pipe! */
        close(apipe[1]);
        close(0);
        dup(apipe[0]);
        close(apipe[0]);
    }
}
```

Real Examples II

```
    execvp(cmd[1][0], cmd[1]);  
    perror("Child_exec::"); exit(1);  
}  
else { /* Parent is "ps -aux" */  
/* We want the stdout connected to pipe */  
    close(apipe[0]);  
    close(1);  
    dup(apipe[1]);  
    close(apipe[1]);  
  
    execvp(cmd[0][0], cmd[0]);  
    perror("Parent_exec::"); exit(1);  
}  
}
```

Real Examples I

```
#include <fcntl.h>
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>
#include <stdlib.h>
/* We want to execute "sort < filecomm1.c | grep fprintf | wc" */
main() {
    int apipe[2];
    int isParent;
    char *cmd[3][3];
    int i, lastChild, fd, saveStdout;
    cmd[0][0] = "sort"; cmd[0][1] = NULL;
    cmd[1][0] = "grep"; cmd[1][1] = "fprintf"; cmd[1][2] = NULL;
    cmd[2][0] = "wc"; cmd[2][1] = NULL;
    saveStdout = dup(1);
    for(i = 2; i >= 0; i--) {
        pipe(apipe);
        isParent = fork();

        if (!isParent) {
            close(apipe[1]);
            close(0);
            if(i != 0) { dup(apipe[0]); }
        }
    }
}
```

Real Examples II

```
    if (i == 0) {
        fd = open("filecomm1.c", O_RDONLY);
        dup(fd);
    }
    close(apipe[0]);

    execvp(cmd[i][0], cmd[i]);
    exit(1);
}
else {
    if (i==2) lastChild = isParent;
    close(apipe[0]);
    close(1);
    if (i!=0) { dup(apipe[1]); }
    close(apipe[1]);
    if (i==0){
        dup2(saveStdout, 1);
        waitpid(lastChild, NULL, 0);
    }
}
}
}
```