

# CS 3411 Systems Programming

Department of Computer Science  
Michigan Technological University

## Sockets

# Today's Topics

- ▶ New Way of Communicating Between Processes
- ▶ Sockets

# "Standard" Unix Processes/IPC

- ▶ IPC stands for Inter-Process Communication in this context
- ▶ Until now, we did communication through:
  - ▶ File system: File descriptors, `read()`/`write()` and pipes
  - ▶ IPC construct is shared through normal process hierarchy inheritance rules, pipes created through `pipe()` are nameless
  - ▶ Totally reliable *byte stream* between producer and consumer
  - ▶ Ties in to conventional UNIX semantics of process creation and termination

## New Ideas

- ▶ We want to create a generalization of the pipe construct for network-based I/O
- ▶ That means we still want file descriptors and `read()/write()` calls to work
- ▶ We need to take some extra features into consideration:
  - ▶ Network Protocol Stacks
  - ▶ Network Naming Conventions
  - ▶ Requirements of Protocol-Specific Message Passing
- ▶ The BSD and UNIX solution is the `socket()` call. Most concisely, it can be described as a *communication endpoint*.
- ▶ The call returns a file descriptor.
- ▶ `int socket(int domain, int type, int protocol);`

# Communication Domain

- ▶ This basically specifies a *protocol stack*.
- ▶ Some systems contain a richer set of communication domains than others
  - ▶ AF\_UNIX or AF\_LOCAL: The UNIX IPC domain, local to a single machine
  - ▶ AF\_INET: The Internet domain, global in scope
  - ▶ AF\_INET6: The Internet domain, using IPv6
- ▶ Once a domain is specified, we know how to associate a name with the socket
- ▶ As well as knowing the semantics of supported IPC mechanisms

# Unix Domain Sockets (in brief)

- ▶ Let's start with the simpler (but less interesting) case of the AF\_UNIX communication domain
- ▶ The header file <sys/un.h> defines addresses

```
#define UNIX_PATH_MAX 108

struct sockaddr_un {
    unsigned short sun_family; /* AF_UNIX */
    char sun_path[UNIX_PATH_MAX]; /* Pathname */
};
```

- ▶ Some examples of Unix domain sockets can be found under the /dev directory

# Types of Sockets in Unix Domain

- ▶ We'll be looking at two *types* of sockets available in Unix
  - ▶ SOCK\_DGRAM provides *datagram* communication semantics
  - ▶ Only promises best-effort delivery
  - ▶ Unix may discard datagrams in times of buffer congestion
  - ▶ Connectionless!

# Types of Sockets in Unix Domain

- ▶ We'll be looking at two *types* of sockets available in Unix
  - ▶ SOCK\_STREAM implements *virtual circuit* communication semantics
  - ▶ Reliable FIFO point-to-point communications
  - ▶ Appears as a byte stream to applications
  - ▶ This is actually how some later UNIX systems implement pipes!
- ▶ Socket type should be chosen according to the needs of the application, and should be programmed in accordance with well-specified delivery semantics of chosen type.



# Operations on Sockets

- ▶ Binding a name to a socket:

```
int bind(int sockfd, struct sockaddr *my_addr, int addrlen);
```

- ▶ Sending datagram on a socket (asynchronous):

```
int sendto(int s, const void *msg, int len, unsigned int flags,  
           const struct sockaddr *to, int tolen);
```

- ▶ Receiving datagram from a socket (synchronous, blocking):

```
int recvfrom(int s, void *buf, int len, unsigned int flags,  
             struct sockaddr *from, int *fromlen);
```

# Server I

```
#include <errno.h>
#include <strings.h>
#include <stdio.h>
#include <unistd.h>
#include <sys/socket.h>
#include <sys/un.h>

main() {
    short p_len;
    int socket_fd, cc, h_len, fsize, namelen;
    void printsun();
    struct sockaddr_un s_un, from;
    size_t addrlen;

    struct {
        char head;
        u_long body;
        char tail;
    } msg;

    socket_fd = socket (AF_UNIX, SOCK_DGRAM, 0);

    s_un.sun_family = AF_UNIX;
```

## Server II

```
strcpy(s_un.sun_path, "udgram");
addrlen = sizeof(s_un.sun_family) + sizeof(s_un.sun_path);
unlink("udgram");

bind(socket_fd, (struct sockaddr *)&s_un, addrlen);

for(;;) {
    fsize = sizeof(from);
    cc = recvfrom(socket_fd, &msg, sizeof(msg), 0,
        (struct sockaddr *)&from, &fsize);
    printsun(&from, "unix_rdgram: ", "Packet from");
    printf("Got data:: %c%d%c\n", msg.head, msg.body, msg.tail);
    fflush(stdout);
}

void printsun(Sun, s1, s2)
struct sockaddr_un *Sun; char *s1, *s2;
{
    printf("%s%s\n", s1, s2);
    printf("family <%d> addr <%s>\n",
        Sun->sun_family, Sun->sun_path);
}
```

# Client I

```
#include <errno.h>
#include <strings.h>
#include <stdio.h>
#include <unistd.h>
#include <sys/socket.h>
#include <sys/un.h>
```

```
main() {
    int socket_fd, cc;
    //long getpid();
    struct sockaddr_un dest;
```

```
    struct {
        char head;
        u_long body;
        char tail;
    } msgbuf;
```

```
    socket_fd = socket (AF_UNIX, SOCK_DGRAM, 0);
    dest.sun_family = AF_UNIX;
    strcpy(dest.sun_path, "udgram");
```

```
    msgbuf.head = '<';
```

## Client II

```
msgbuf.body = (u_long) getpid();  
msgbuf.tail = '>';  
  
cc = sendto(socket_fd, &msgbuf, sizeof(msgbuf), 0,  
            (struct sockaddr *)&dest, sizeof(dest));  
}
```

# Sockets and the Internet (IPv4)

- ▶ `AF_INET` communication domain
- ▶ `SOCK_DGRAM` - Same as before! **UDP/IP**
- ▶ `SOCK_STREAM` - Same as before! **TCP/IP**
- ▶ We need a way to associate names with sockets to be able to do network I/O through a socket file descriptor

# Sockets and the Internet (IPv4)

- ▶ The header file `<netinet/in.h>` defines a 32-bit for an Internet host.
- ▶ This actually identifies a specific network interface on a specific system on the Internet.
- ▶ It's represented by a 32 bit unsigned number

```
struct in_addr {  
    __u32 s_addr;  
}
```