

CS 3411 Systems Programming

Department of Computer Science
Michigan Technological University

Sockets

Today's Topics

- ▶ New Way of Communicating Between Processes
- ▶ Sockets

Sockets and the Internet (IPv4)

- ▶ AF_INET communication domain
- ▶ SOCK_DGRAM - Provides *datagram* semantics, only promises best-effort delivery! **UDP/IP**
- ▶ SOCK_STREAM - Provides a FIFO type point-to-point communication system! **TCP/IP**
- ▶ We need a way to associate names with sockets to be able to do network I/O through a socket file descriptor

Sockets and the Internet (IPv4)

- ▶ The header file `<netinet/in.h>` defines a 32-bit for an Internet host.
- ▶ This actually identifies a specific network interface on a specific system on the Internet.
- ▶ It's represented by a 32 bit unsigned number

```
struct in_addr {  
    __u32 s_addr;  
}
```

- ▶ The addresses are usually represented by dotted decimal notation.

Representing the Address in C

- ▶ In header file <netinet/in.h>

```
#define __SOCK_SIZE__ 16 /* sizeof(struct sockaddr) */

struct sockaddr_in {
    short int sin_family; /* Address family */
    unsigned short int sin_port; /* Port number */
    struct in_addr sin_addr; /* Internet address */
    /* Pad to size of 'struct sockaddr'. */
    unsigned char __pad[__SOCK_SIZE__ - sizeof(short int) -
        sizeof(unsigned short int) - sizeof(struct in_addr)];
};
```

- ▶ Declare and/or allocate instance of struct sockaddr_in whenever you need to specify a full address on the Internet
- ▶ A port is an Internet communication endpoint associated with an application. (*host,port*) defines an Internet address.
- ▶ Ports in range [0,1023] reserved for root; others available to ordinary users. (See RFC 1700)

Usual Ports for Services

- ▶ FTP uses 20 and 21
- ▶ SSH uses 22
- ▶ Telnet uses 23
- ▶ HTTP uses 80, commonly
- ▶ HTTPS uses 443
- ▶ Check `/etc/services` to see what "well-known" ports are

Translating Host Names into IP Address(es)

- ▶ Library function to map symbolic host name into IP address(es):

```
#include <netdb.h>
```

```
struct hostent *gethostbyname(const char *name);
```

```
void herror(const char *s);
```

- ▶ The hostent data structure:

```
struct hostent {  
    char *h_name;           /* official name of host */  
    char **h_aliases;       /* alias list */  
    int h_addrtype;         /* host address type */  
    int h_length;           /* length of address */  
    char **h_addr_list;     /* list of addresses */  
}
```

```
#define h_addr h_addr_list[0] /* for backward compatibility */
```

Translating Host Names into IP Address(es)

- ▶ If we have a dotted decimal string and we want to convert it into an address we can use, the above function is useful.
- ▶ Also see `man inet` for more functions!

getaddrs.c Example

```
#include <netdb.h>
#include <stdio.h>
#include <stdlib.h>
#include <sys/socket.h>
#include <arpa/inet.h>
#include <netinet/in.h>
main(argc, argv)
int argc; char **argv;
{
    struct hostent *entry; char **next;
    struct in_addr address, **addrptr;
    entry = gethostbyname(argv[1]);
    if (!entry) { perror("lookup error"); exit(1); }
    printf("Official name->%s\n", entry->h_name);
    if (entry->h_aliases[0]) {
        printf("Aliases->\n");
        for (next = entry->h_aliases; *next; next++)
            printf("%s\n", *next);
    }
    printf("IP Addresses:\n");
    for (addrptr=(struct in_addr **) entry->h_addr_list;
        *addrptr; addrptr++)
        printf("%s\n", inet_ntoa(**addrptr));
}
```

Can Also Get Symbolic Name from IP Address

- ▶ There's also an inverse function (we know IP address, want symbolic name)

```
#include <netdb.h>
```

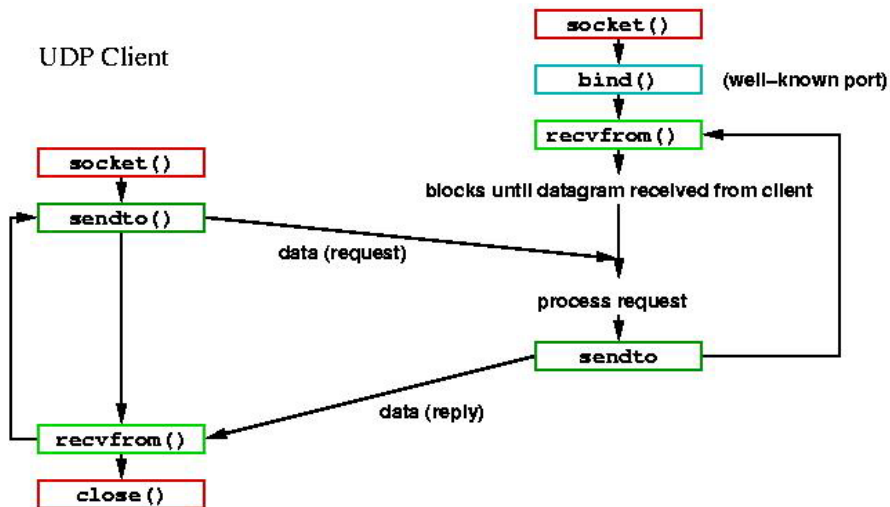
```
struct hostent *gethostbyaddr(const char *addr,  
                               int len, int type);
```

gethost.c Example

```
#include <netdb.h>
#include <stdio.h>
#include <stdlib.h>
#include <sys/socket.h>
#include <arpa/inet.h>
#include <netinet/in.h>
main(argc, argv)
int argc; char **argv;
{
    struct hostent *entry; char **next;
    struct in_addr address, **addrptr;
    inet_aton(argv[1], &address);
    entry = gethostbyaddr((char *)&address, sizeof(address),
        AF_INET);
    if (!entry) { perror("lookup error"); exit(1); }
    printf("Official name->%s\n", entry->h_name);
    if (entry->h_aliases[0]) {
        printf("Aliases->\n");
        for (next = entry->h_aliases; *next; next++)
            printf("%s\n", *next);
    }
    printf("IP Addresses:\n");
    for (addrptr=(struct in_addr **) entry->h_addr_list;
        *addrptr; addrptr++)
        printf("%s\n", inet_ntoa(**addrptr));
}
```

UDP Server

UDP Client



recv_upd.c: UDP/IP Server I

```
#include <netdb.h>
#include <stdio.h>
#include <stdlib.h>
#include <sys/socket.h>
#include <arpa/inet.h>
#include <netinet/in.h>
#include <strings.h>

void printsin(s_in, s1, s2)
struct sockaddr_in *s_in; char *s1, *s2;
{
    printf ("Program: %s\n%s", s1, s2);
    printf ("(%d,%d)\n", s_in->sin_addr.s_addr, s_in->sin_port);
}

main()
{
    short p_len;
    int socket_fd, cc, h_len, fsize, namelen;
    struct sockaddr_in s_in, from;
    struct { char head; u_long body; char tail;} msg;

    socket_fd = socket (AF_INET, SOCK_DGRAM, 0);
    /* You must do this just in case */
```

recv_upd.c: UDP/IP Server II

```
bzero((char *) &s_in, sizeof(s_in));

s_in.sin_family = (short)AF_INET;
s_in.sin_addr.s_addr = htonl(INADDR_ANY); /* WILDCARD */
s_in.sin_port = htons((u_short)0x3333);
printsln( &s_in, "RECV_UDP", "Local_socket is:");
fflush(stdout);
bind(socket_fd, (struct sockaddr *)&s_in, sizeof(s_in));
for(;;) {
    fsize = sizeof(from);
    cc = recvfrom(socket_fd, &msg, sizeof(msg), 0,
        (struct sockaddr *)&from, &fsize);
    printsln( &from, "recv_udp:", "Packet from:");
    printf("Got data:: %c%d%c\n", msg.head,
        ntohs(msg.body), msg.tail);
    fflush(stdout);
}
}
```

send_upd.c: UDP/IP Client I

```
#include <netdb.h>
#include <stdio.h>
#include <stdlib.h>
#include <sys/socket.h>
#include <arpa/inet.h>
#include <netinet/in.h>
#include <strings.h>
main(argc, argv)
int argc; char **argv;
{
    int socket_fd;
    struct sockaddr_in dest;
    struct hostent *hostptr;
    struct { char head; u_long body; char tail; } msgbuf;
    socket_fd = socket (AF_INET, SOCK_DGRAM, 0);
    /* You must do this just in case */
    bzero((char *) &dest, sizeof(dest));

    hostptr = gethostbyname(argv[1]);
    dest.sin_family = (short) AF_INET;
    bcopy(hostptr->h_addr, (char *)&dest.sin_addr,
        hostptr->h_length);
    dest.sin_port = htons((u_short)0x3333);
```

send_upd.c: UDP/IP Client II

```
msgbuf.head = '<';  
msgbuf.body = htonl(getpid()); /* IMPORTANT! */  
msgbuf.tail = '>';  
sendto(socket_fd, &msgbuf, sizeof(msgbuf), 0,  
        (struct sockaddr *)&dest, sizeof(dest));  
}
```


Similarities and Differences

- ▶ Note that there are striking similarities between Unix datagram programs and Internet datagram programs
- ▶ We do need to do extra work for Internet programs
- ▶ Socket creation parameters are trivially different
- ▶ Naming conventions are significantly
- ▶ The underlying implementation is completely different! But hidden from the programmers
- ▶ Practical note: You can always test and develop network programs on "localhost" (127.0.0.1). The implementation should be smart enough to *NOT* send the packets over the network (instead just pass it from output buffer to input buffer)