

# CS 3411 Systems Programming

Department of Computer Science  
Michigan Technological University

## Sockets

# Today's Topics

- ▶ TCP Connections
- ▶ More details about the inet domain

# Fancier UDP Application

- ▶ Let's make our application a little bit fancier
- ▶ We want it to stop if we don't receive a datagram after 1 minute
- ▶ We also want it to stop if anything is written in stdin
- ▶ The problem: Having a read/recv call hanging from several descriptors at once
- ▶ Solution: New kernel call!

```
#include <sys/time.h>
#include <sys/types.h>
#include <unistd.h>
```

```
int select(int nfds, fd_set *readfds, fd_set *writefds,
           fd_set *exceptfds, struct timeval *timeout);
```

```
void FD_CLR(int fd, fd_set *set);
int  FD_ISSET(int fd, fd_set *set);
void FD_SET(int fd, fd_set *set);
void FD_ZERO(fd_set *set);
```

# Fancy UDP Server I

```
#include <sys/select.h>
#include <netdb.h>
#include <stdio.h>
#include <stdlib.h>
#include <sys/socket.h>
#include <arpa/inet.h>
#include <netinet/in.h>
#include <strings.h>
void printsin(s_in, s1, s2)
struct sockaddr_in *s_in; char *s1, *s2;
{
    printf ("Program: %s\n%s", s1, s2);
    printf ("(%d,%d)\n", s_in->sin_addr.s_addr, s_in->sin_port);
}

main() {
    int socket_fd, cc, h_len, fsize, namelen, hits;
    fd_set mask;
    struct timeval timeout;
    struct sockaddr_in s_in, from;

    struct { char head; u_long body; char tail; } msg;
```

## Fancy UDP Server II

```
socket_fd = socket(AF_INET, SOCK_DGRAM, 0);

bzero((char *) &s_in, sizeof(s_in));
s_in.sin_family = (short) AF_INET;
s_in.sin_addr.s_addr = htons(INADDR_ANY);
s_in.sin_port = htons((u_short)0x3333);
bind(socket_fd, (struct sockaddr *)&s_in, sizeof(s_in));

for(;;) {
    fsize = sizeof(from);
    FD_ZERO(&mask); FD_SET(0, &mask); FD_SET(socket_fd, &mask);
    timeout.tv_sec = 60; timeout.tv_usec = 0;

    if ((hits = select(socket_fd+1, &mask, (fd_set *)0,
                      (fd_set *)0, &timeout)) < 0) {
        perror("recv_udp: select"); exit(1);
    }

    if( (hits == 0) || ((hits > 0) && (FD_ISSET(0, &mask)))) {
        printf("Shutting down\n"); exit(0);
    }

    cc = recvfrom(socket_fd, &msg, sizeof(msg), 0,
                  (struct sockaddr *)&from, &fsize);
```

# Fancy UDP Server III

```
    printsin(&from, "recv_udp: ", "Packet from:");  
    printf("Got data:: %c%d%c\n", msg.head,  
          ntohl(msg.body), msg.tail);  
    fflush(stdout);  
}  
}
```

# Internet Virtual Circuits (TCP/IP)

- ▶ We want to extend the pipe abstraction onto the Internet
- ▶ This means a reliable byte stream with no visible packets/messages!
- ▶ Also implies a point to point connection. In the entire Internet, the tuple
  - ▶  $((\text{host1}, \text{port1}), (\text{host2}, \text{port2}))$is unique
- ▶ It should be noted that the cost for setting up, maintaining and tearing down the connection has a non-trivial cost
- ▶ TCP/IP implies a client-server model!

# Making TCP/IP Connections!

- ▶ New kernel calls for TCP/IP

- ▶ To mark a socket as being capable of accepting TCP/IP connections, the server calls:

```
#include <sys/socket.h>
int listen(int s, int backlog);
```

- ▶ To mark a socket as being capable of accepting TCP/IP connections, the server calls:

```
#include <sys/socket.h>
int listen(int s, int backlog);
```

- ▶ To accept a TCP/IP connection on a listening socket, server can then do:

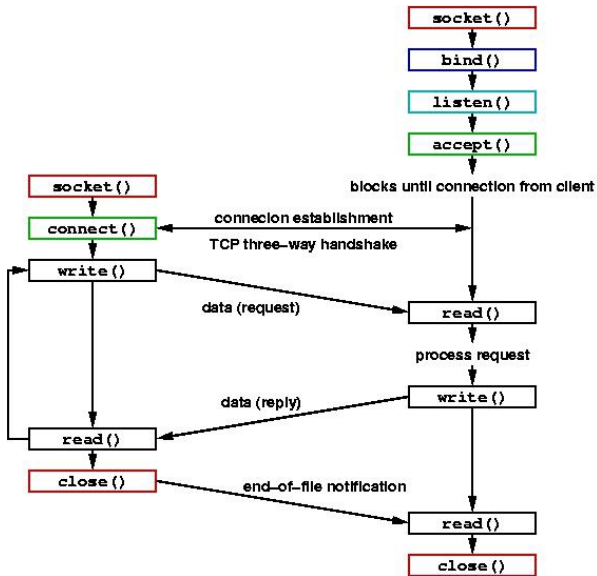
```
#include <sys/types.h>
#include <sys/socket.h>
int accept(int s, struct sockaddr *addr, int *addrlen);
```

- ▶ To initiate a connection to a server, client does:

```
#include <sys/types.h>
#include <sys/socket.h>
int connect(int sockfd, struct sockaddr *serv_addr,
            int addrlen);
```

- ▶ Then we can use read and write to pass data over the circuit!





# TCP/IP Server I

```
#include <sys/socket.h>
#include <sys/types.h>
#include <arpa/inet.h>
#include <netinet/in.h>
#include <strings.h>
#include <netdb.h>
#include <stdlib.h>
#include <stdio.h>
void printsin(s_in, s1, s2)
struct sockaddr_in *s_in; char *s1, *s2;
{
    printf ("Program: %s\n%s", s1, s2);
    printf ("(%d,%d)\n", s_in->sin_addr.s_addr, s_in->sin_port);
}
main() {
    int listener, conn, length; char ch;
    struct sockaddr_in s1, s2;

    listener = socket( AF_INET, SOCK_STREAM, 0 );

    bzero((char *) &s1, sizeof(s1));
    s1.sin_family = (short) AF_INET;
    s1.sin_addr.s_addr = htonl(INADDR_ANY);
```

# TCP/IP Server II

```
s1.sin_port = htons(0); /* bind() will gimme unique port. */
bind(listener, (struct sockaddr *)&s1, sizeof(s1));
length = sizeof(s1);
getsockname(listener, (struct sockaddr *)&s1, &length);
printf("RSTREAM:: assigned port number %d\n",
       ntohs(s1.sin_port));

listen(listener, 1);
length = sizeof(s2);
conn=accept(listener, (struct sockaddr *)&s2, &length);
printsin(&s2, "RSTREAM::", "accepted connection from");
printf("\n\nRSTREAM:: data from stream:\n");
while ( read(conn, &ch, 1) == 1) putchar(ch);
putchar('\n');
}
```

# TCP/IP Client I

```
#include <sys/socket.h>
#include <sys/types.h>
#include <arpa/inet.h>
#include <netinet/in.h>
#include <strings.h>
#include <netdb.h>
#include <stdlib.h>
char msg[] = { "the quick brown fox jumps over the lazy dog" };
main(int argc, char **argv) {
    char *remhost; u_short remport;
    int sock, left, num, put;
    struct sockaddr_in remote;
    struct hostent *h;

    remhost = argv[1]; remport = atoi(argv[2]);

    sock = socket(AF_INET, SOCK_STREAM, 0);

    bzero((char *) &remote, sizeof(remote));
    remote.sin_family = (short) AF_INET;
    h = gethostbyname(remhost);
    bcopy((char *) h->h_addr, (char *) &remote.sin_addr,
          h->h_length);
```

# TCP/IP Client II


```
remote.sin_port = htons(remport);
connect(sock, (struct sockaddr *)&remote, sizeof(remote));

left = sizeof(msg); put = 0;
while(left > 0) {
    if((num = write(sock, msg+put, left)) < 0) {
        perror("inet_wstream: write");
        exit(1);
    }
    else {
        left -= num;
        put += num;
    }
}
}
```


# A Typical TCP/IP Client/Server Server

```
...  
listener = socket (...);  
...  
bind(listener, ...);  
listen(listener, ...);  
...  
while (1) {  
    new = accept(listener, ...);  
    if (fork() == 0) {  
        close(listener);  
        /* read lots of stuff from new */  
    }  
    close(new);  
    while(wait3(status, WNOHANG, NULL)); \  
    /* Can also handle SIGCHLD */  
}
```

# OSI Reference Model



OSI REFERENCE MODEL






INTERNET PROTOCOL  
SUITE

*user process*

---

*kernel*

# Packet Encapsulation





# Ethernet

- ▶ IEEE Standard 802: CSMA/CD, token bus, token ring
- ▶ IEEE Standard 802.3: CSMA/CD
  - ▶ Ethernet is implementation of 802.3
- ▶ Typically 6 byte address: 00:62:97:B8:C9:4A
- ▶ ARP: protocol for mapping IPv4 address the hardware address
- ▶ RARP: protocol for mapping hardware address to IPv4 address
- ▶ Ethernet address outermost
- ▶ Routing protocols determine correct IP address on a hop-by-hop basis
- ▶ ARP maps IP address to ethernet address to transmit message across next hop
- ▶ man arp and man traceroute